

Discriminative models

Greg Tsagkatakis

CSD - UOC

ICS - FORTH

Types of problems



Regression



What will be the temperature tomorrow?

84°



Fahrenheit

Classification



Will it be hot or cold tomorrow?

COLD

HOT



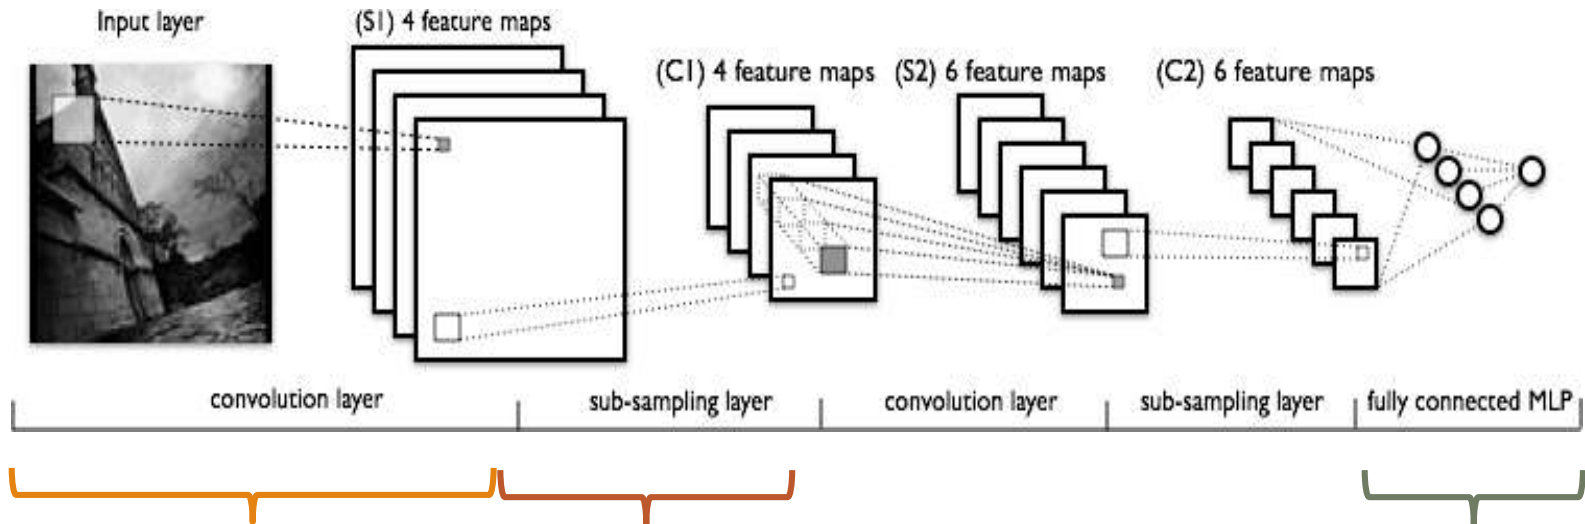
Fahrenheit

State-of-the-art (since 2015)

Deep Learning (DL)

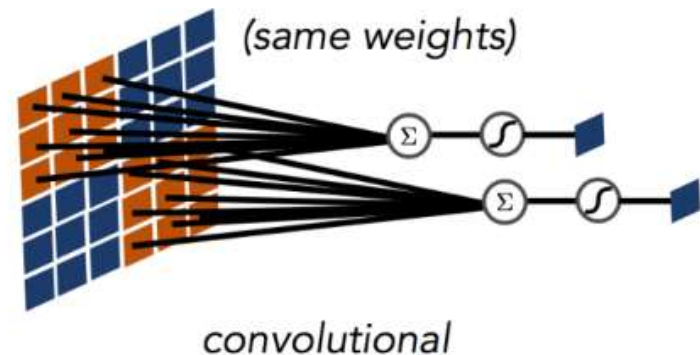
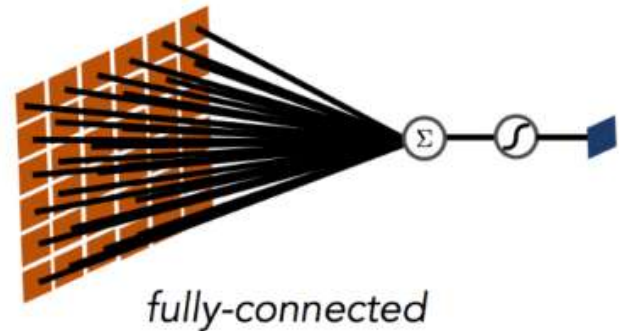
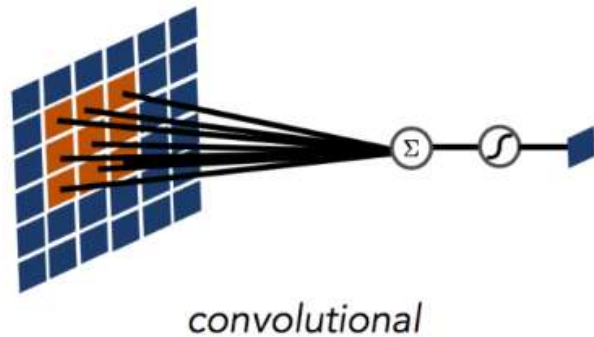
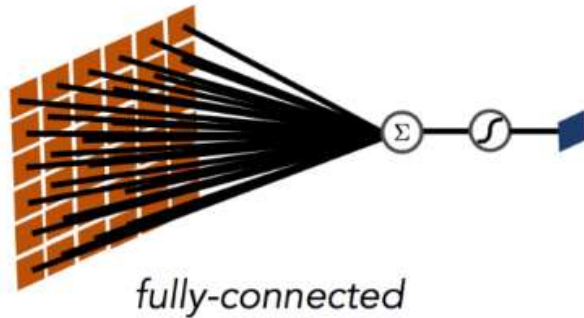
- Convolutional Neural Networks (CNN) \leftrightarrow Images
- Recurrent Neural Networks (RNN) \leftrightarrow Audio

Convolutional Neural Networks



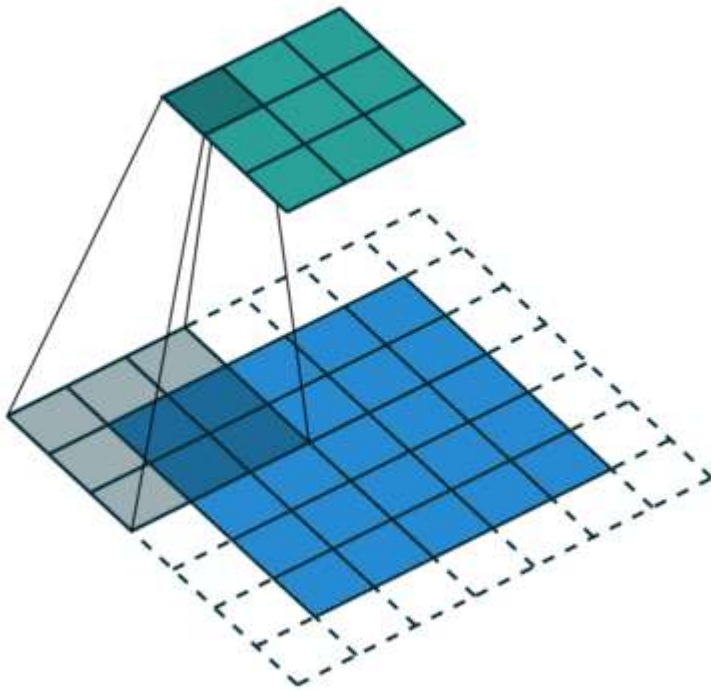
(Convolution + Subsampling) + () ... + Fully Connected

Convolutional Neural Networks



Tutorial: Introduction to convolutional neural networks (CNNs)
https://github.com/langnico/DL_tutorial_RS/tree/master

Convolution operator



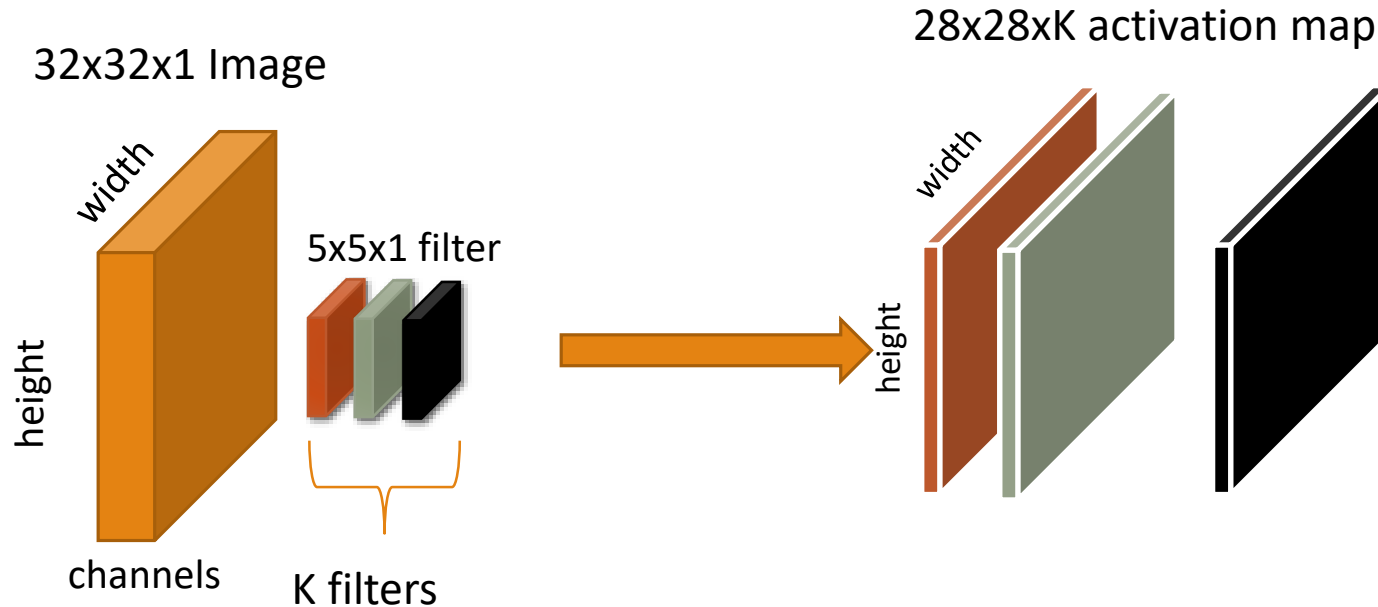
1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolutional Layers



$$\begin{aligned}(I * K)_{ij} &= \sum_{m=0}^{k_1-1} \sum_{n=0}^{k_2-1} I(i-m, j-n) K(m, n) \\ &= \sum_{m=0}^{k_1-1} \sum_{n=0}^{k_2-1} I(i+m, j+n) K(-m, -n)\end{aligned}$$

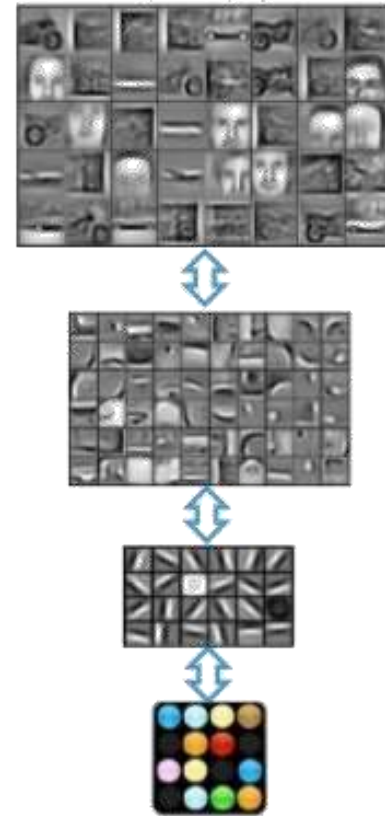
Convolutional Layers

Characteristics

- Hierarchical features
- Location invariance

Parameters

- Number of filters (32,64...)
- Filter size (3x3, 5x5)
- Stride (1)
- Padding (2,4)



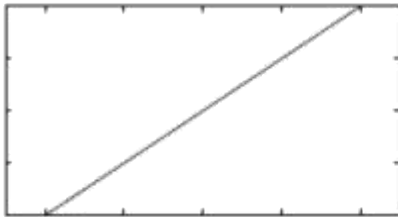
“Machine Learning and AI for Brain Simulations” –
Andrew Ng Talk, UCLA, 2012

Activation Layer

Introduction of non-linearity

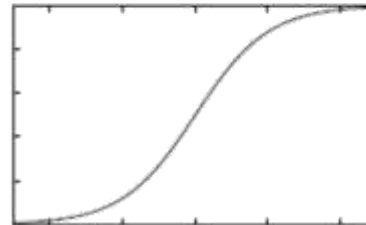
- Brain: thresholding -> spike trains

Identity (Linear)



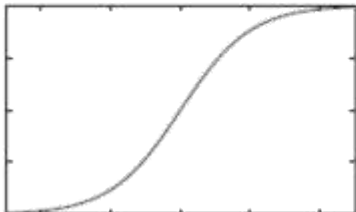
$$\text{identity}(x) = x$$

Sigmoid



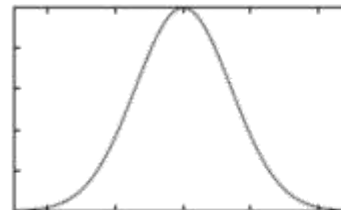
$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

Tanh (Hypertangent)



$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Gaussian



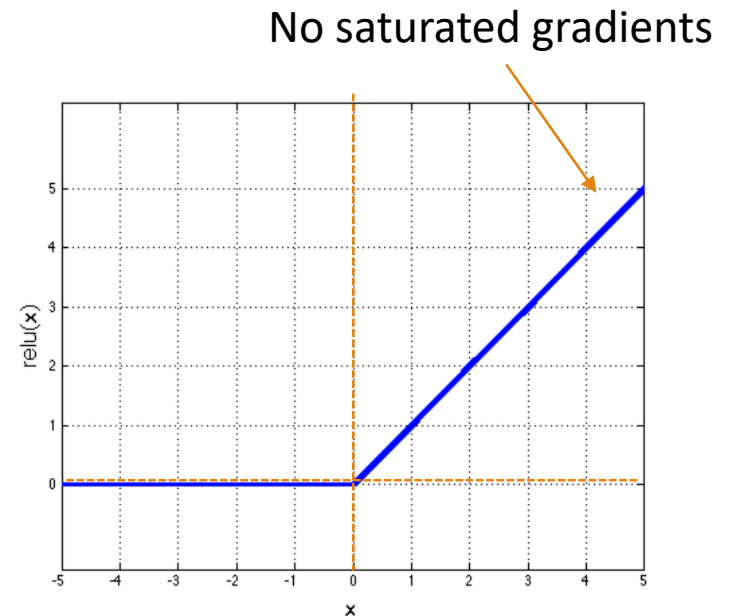
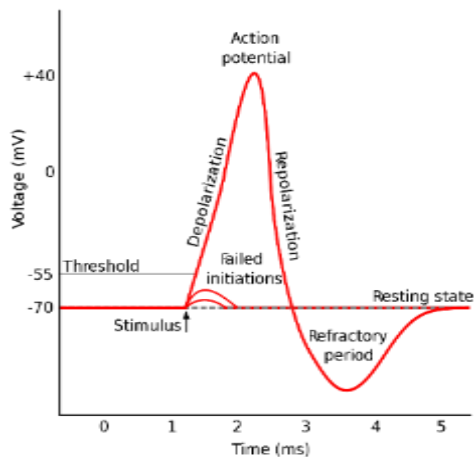
$$\text{gaussian}(x) = e^{-x^2/\sigma^2}$$

Activation Layer

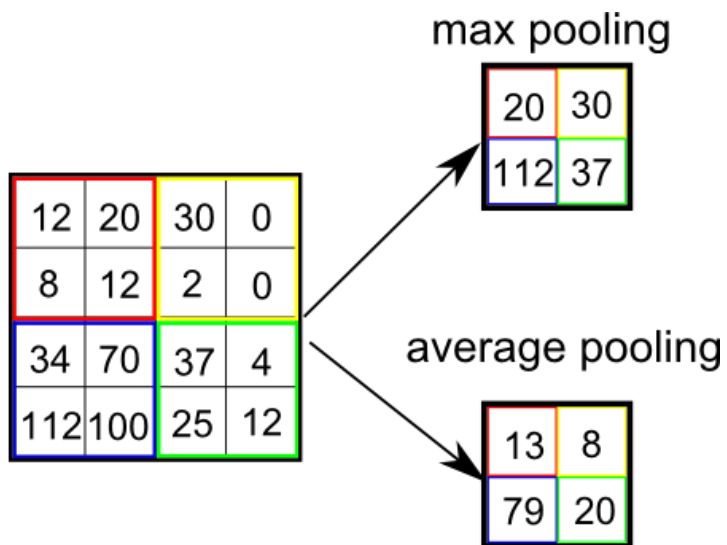
ReLU: $x = \max(0, x)$

- ✓ Simplifies backprop
- ✓ Makes learning faster
- ✓ Avoids saturation issues
- ✓ ~ non-negativity constraint

(Note: The brain)



Subsampling (pooling) Layers



<-> downsampling

➤ Scale invariance

Parameters

- Type
- Filter Size
- Stride

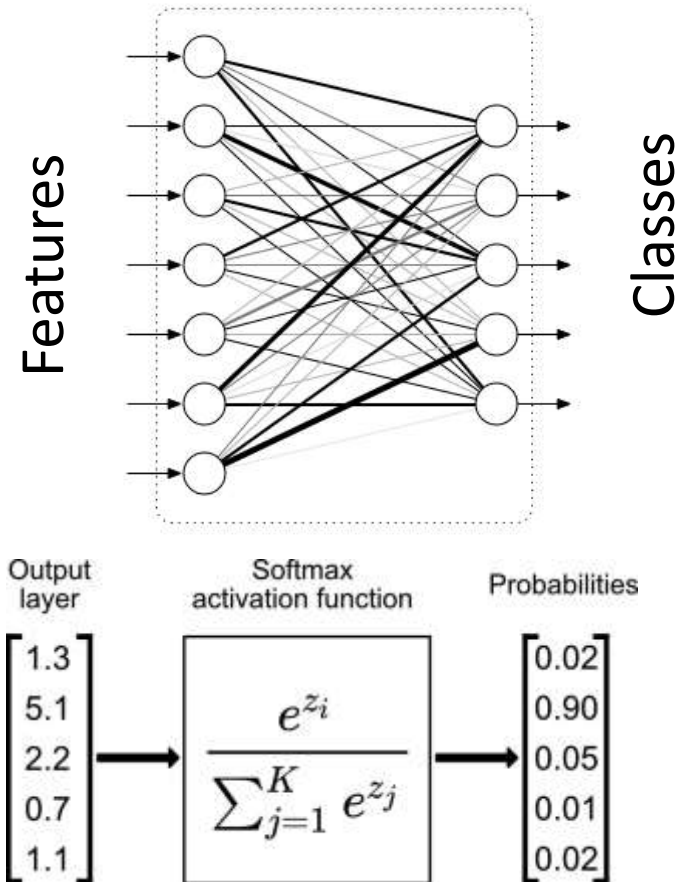
Fully Connected Layers

Full connections to all activations in the previous layer

Typically at the end

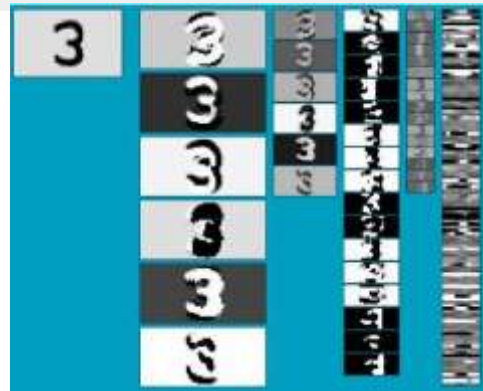
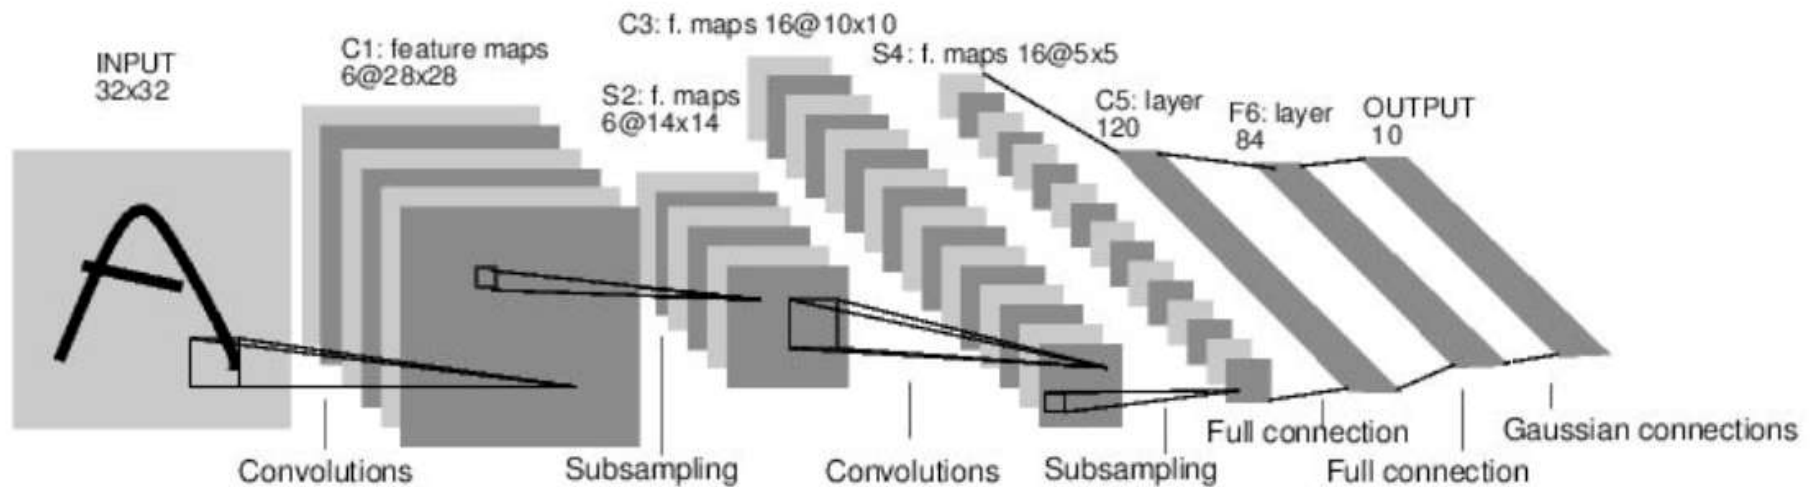
Can be replaced by conv

Softmax on logits

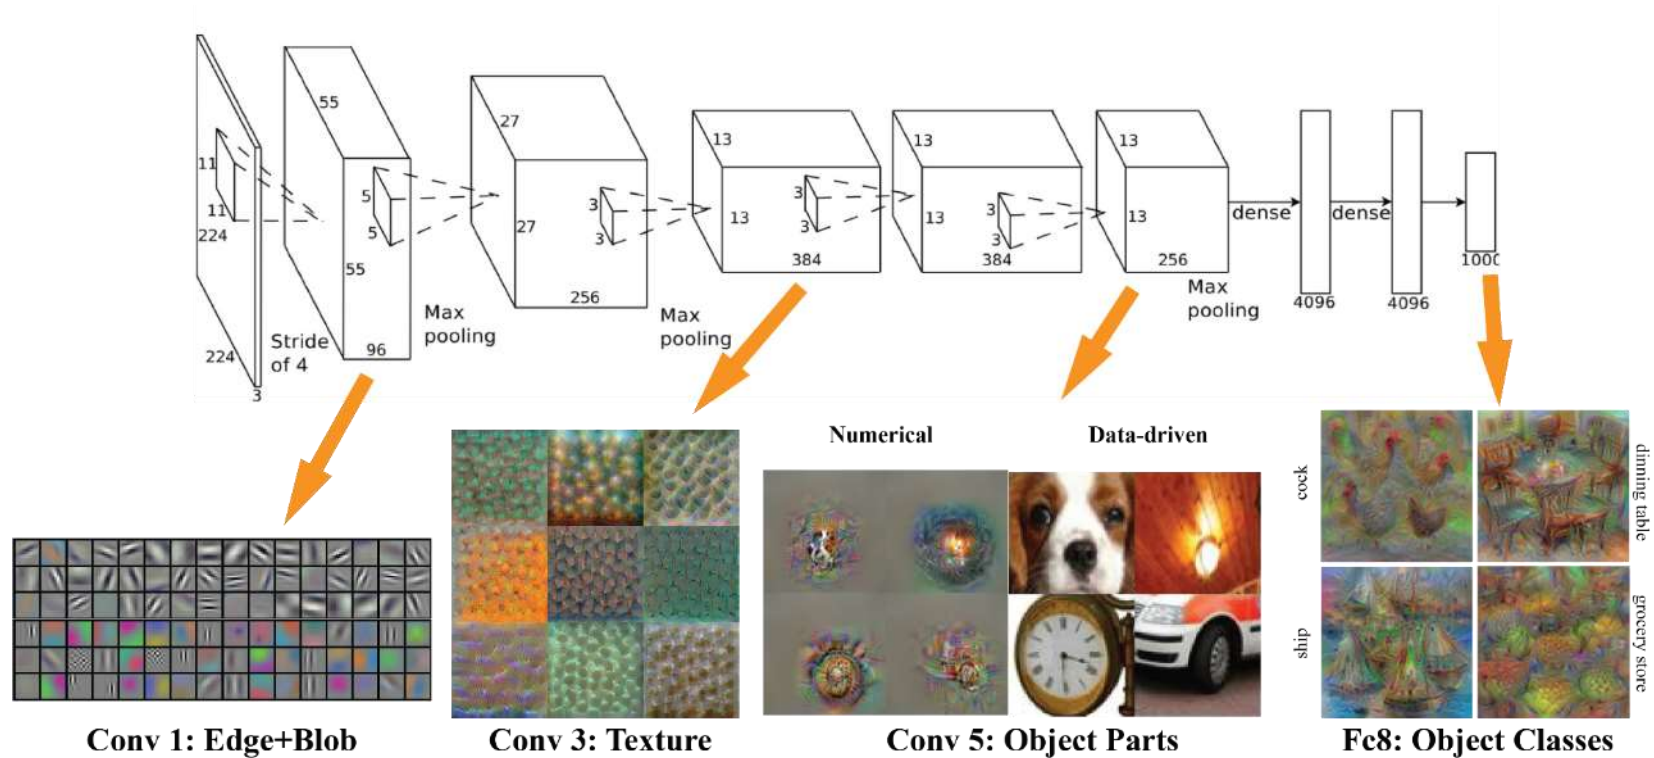


LeNet [1998]

[LeCun et al., 1998]

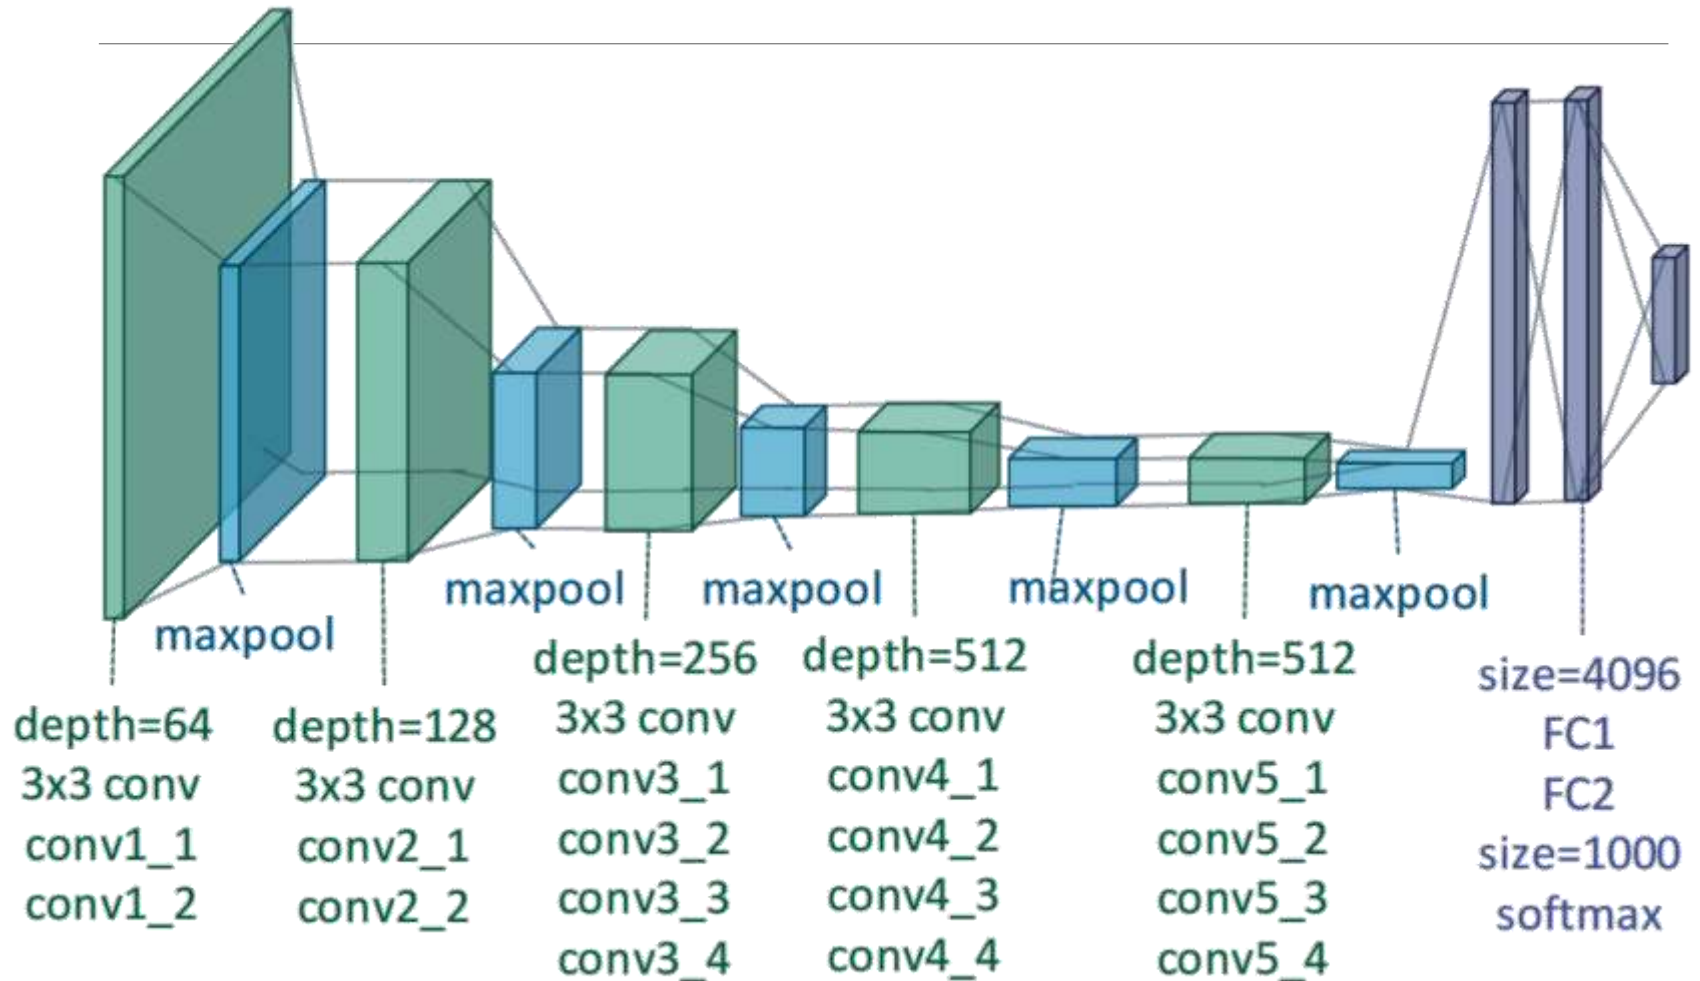


AlexNet [2012]



Alex Krizhevsky, Ilya Sutskever and Geoff Hinton, [ImageNet ILSVRC challenge](http://vision03.csail.mit.edu/cnn_art/data/single_layer.png) in 2012
http://vision03.csail.mit.edu/cnn_art/data/single_layer.png

VGGnet [2014]



K. Simonyan, A. Zisserman Very Deep Convolutional Networks for Large-Scale Image Recognition, arXiv technical report, 2014

VGGnet

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

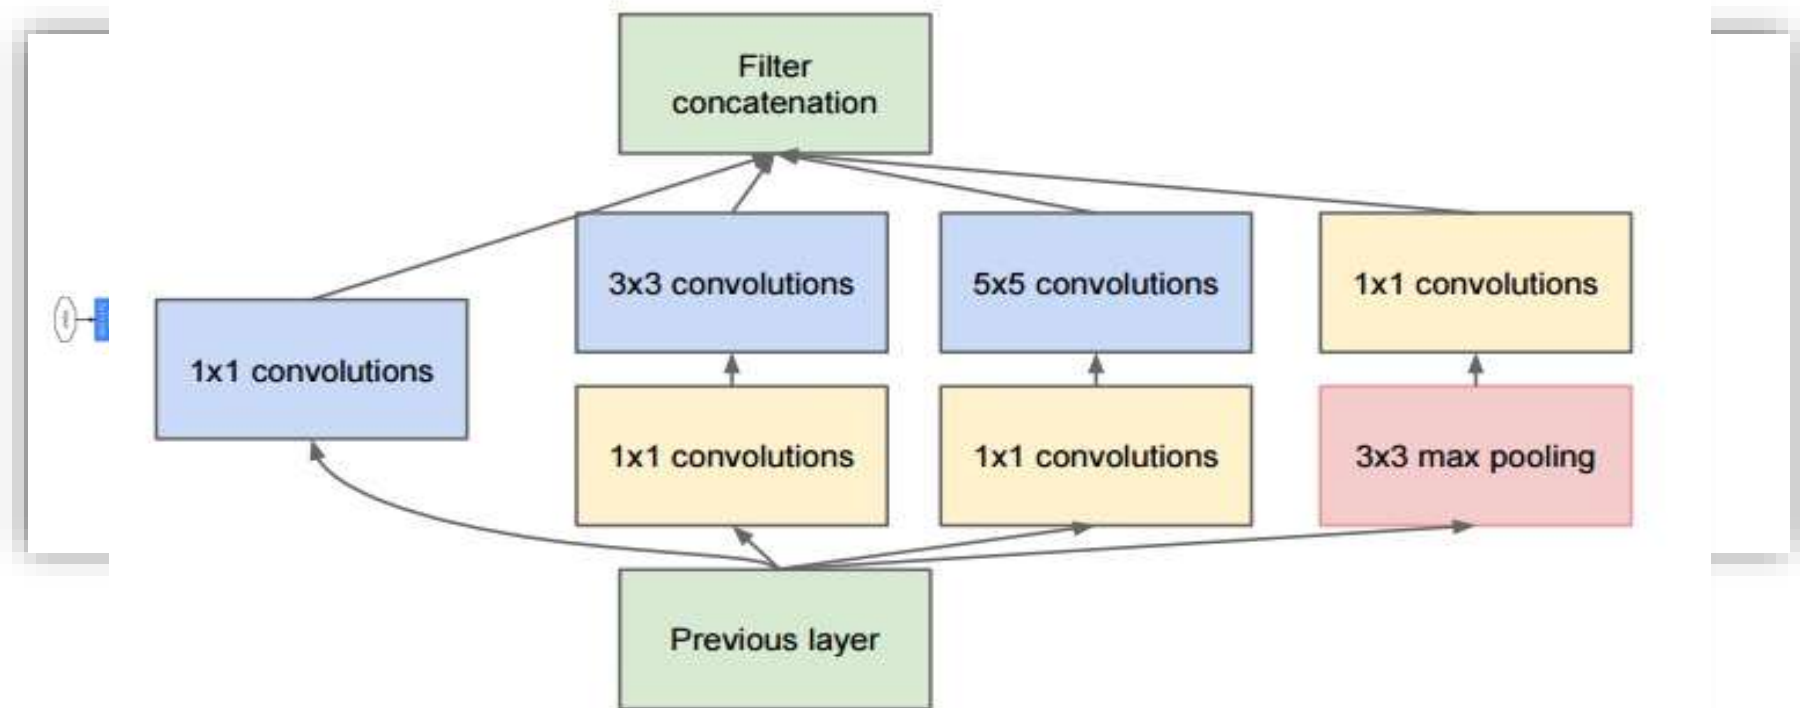
D: VGG16

E: VGG19

All filters are 3x3

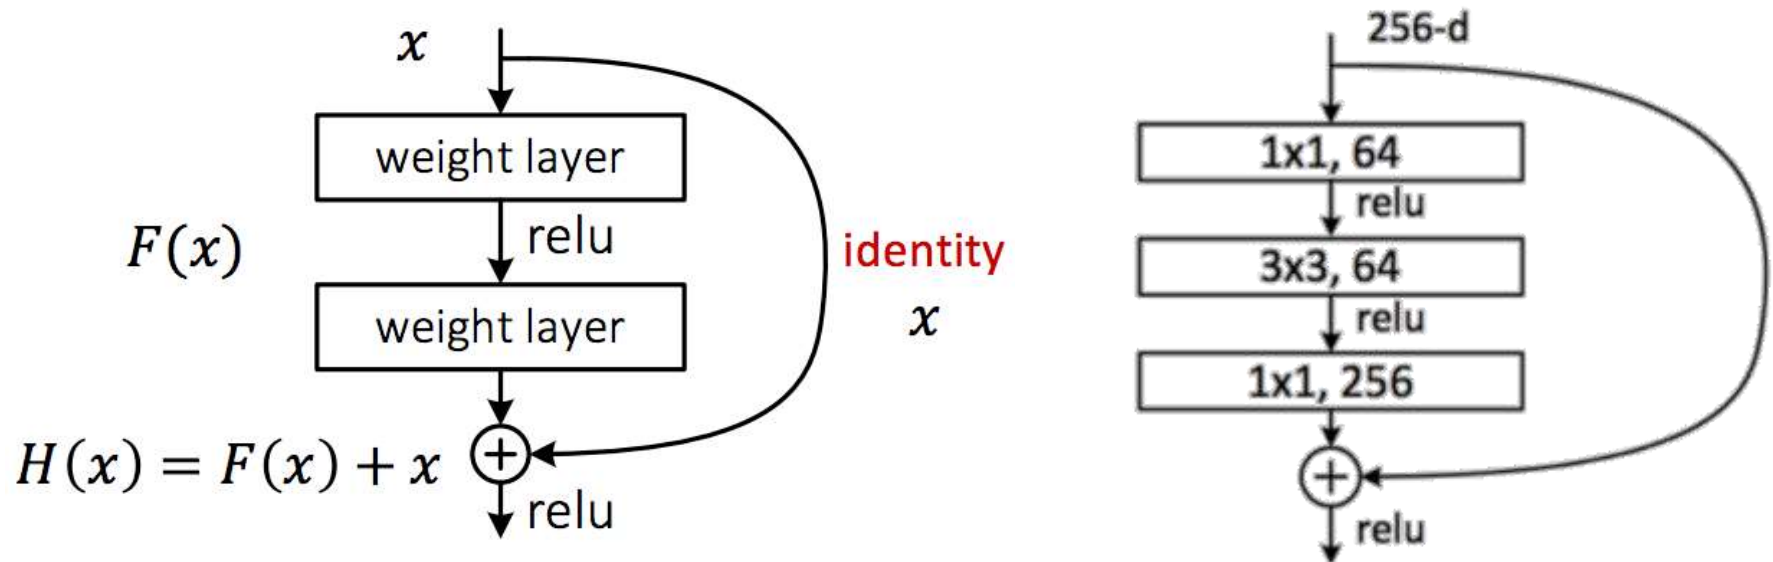
More layers
smaller filters

Inception (GoogLeNet, 2014)



Inception module with dimensionality reduction

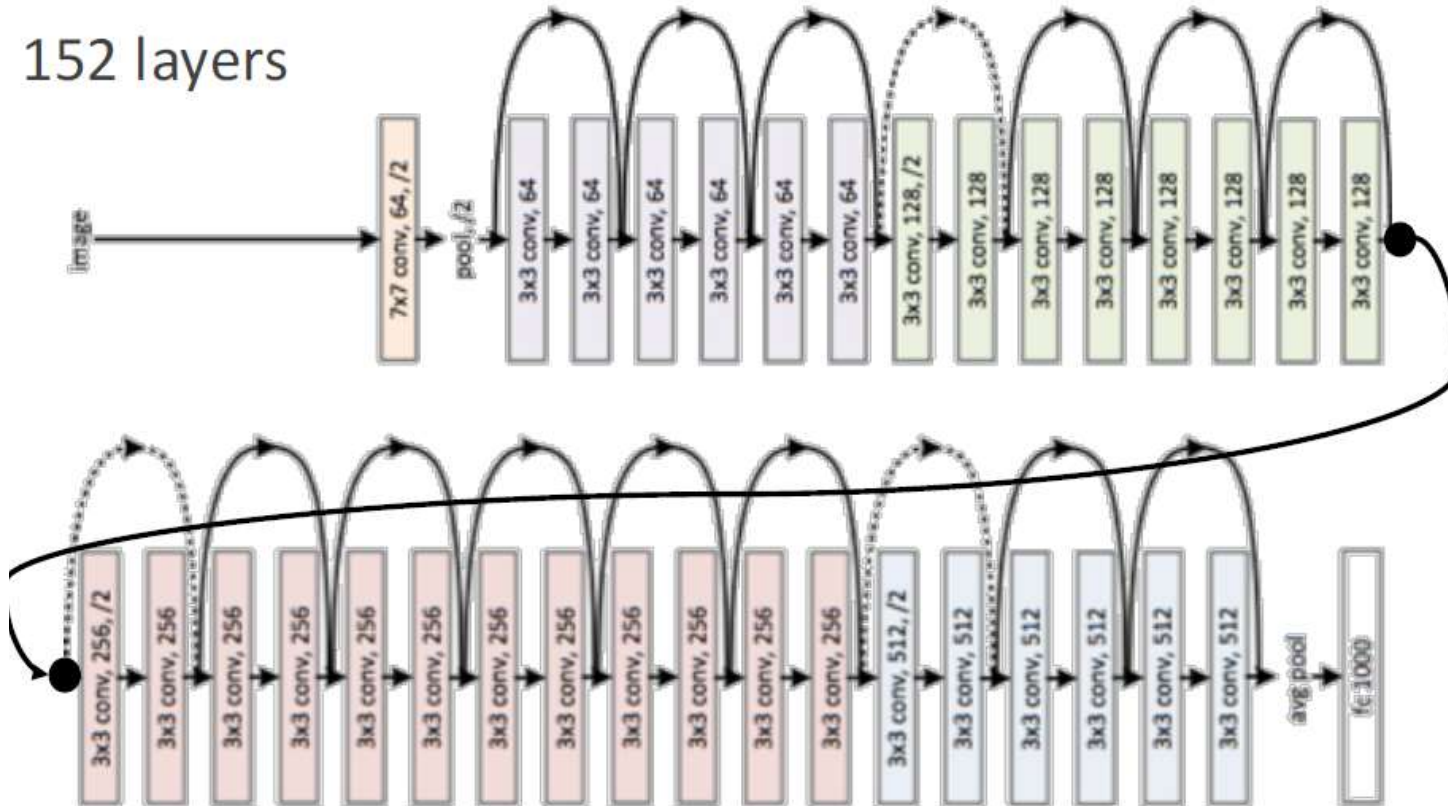
Residuals



ResNet, 2015

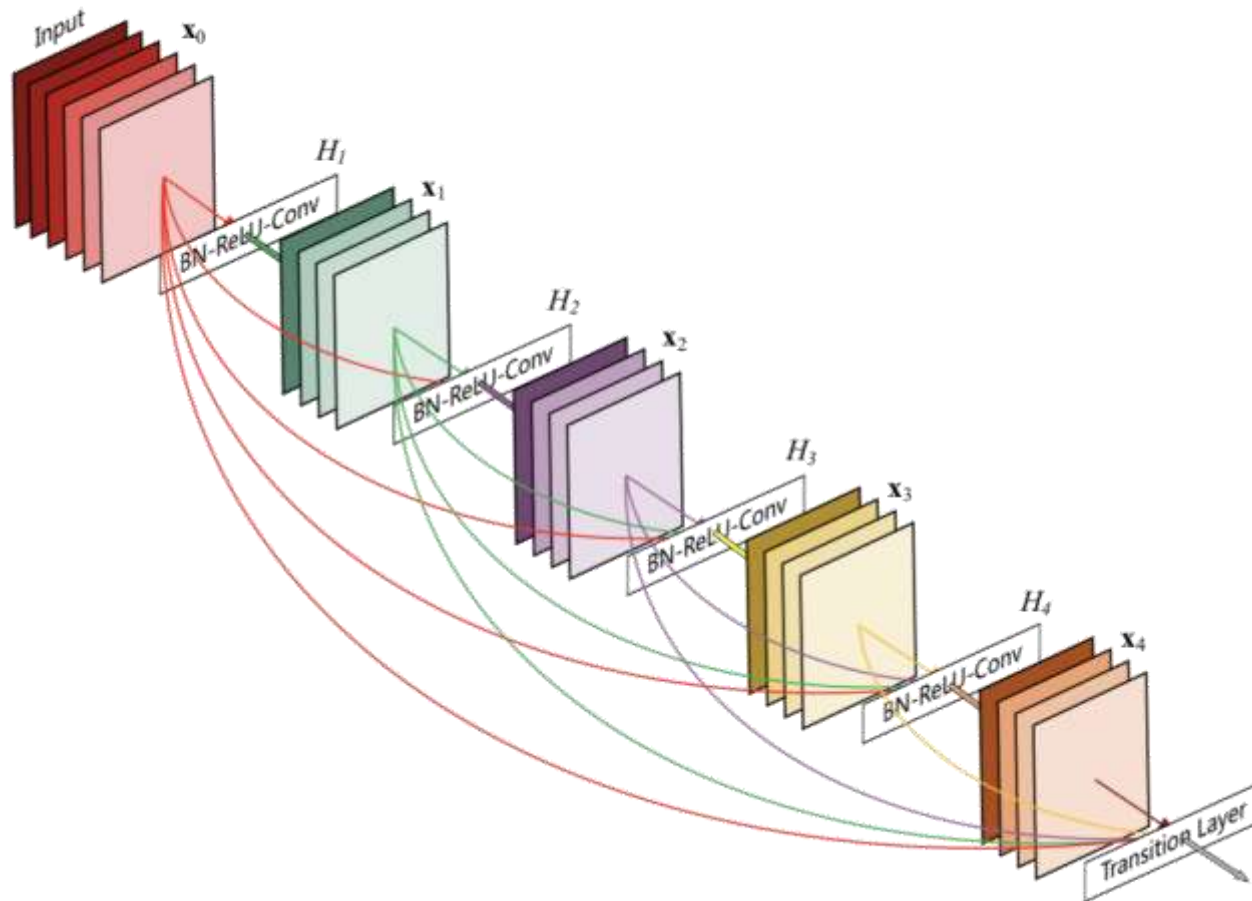
Residual Networks

152 layers



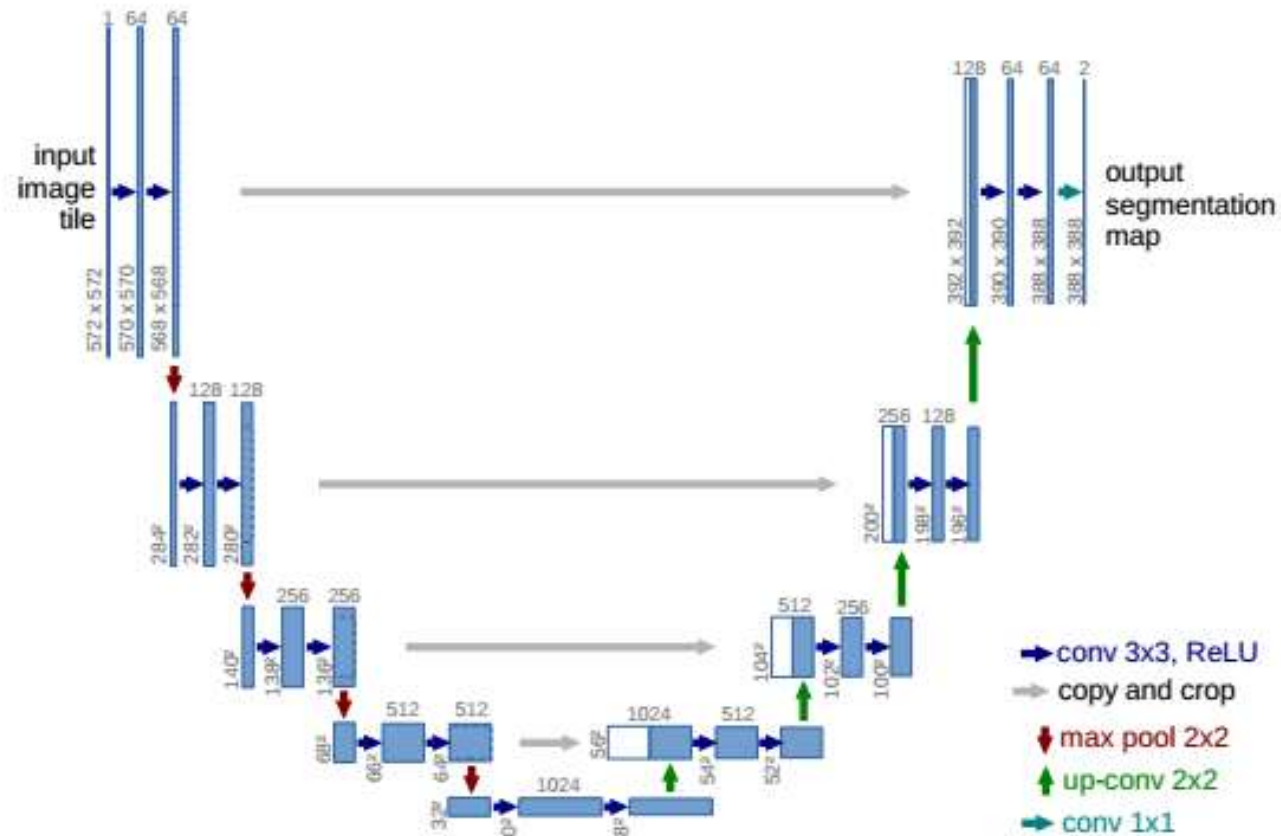
He, Kaiming, et al. "Deep residual learning for image recognition." *IEEE CVPR*. 2016.

DenseNet



Densely Connected Convolutional Networks, 2016

U-NET



Ronneberger, O., Fischer, P., & Brox, T. (2015). U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2015*:

Recurrent Neural Networks

Motivation

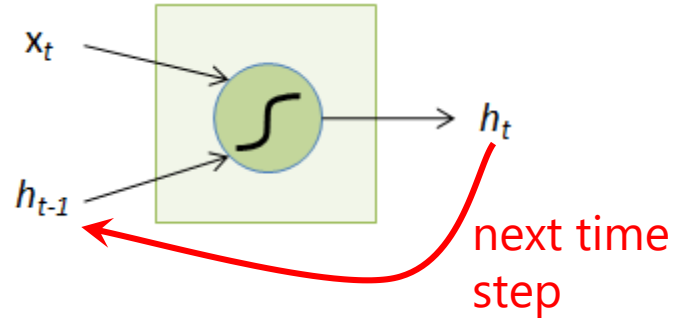
- Feed forward networks accept a fixed-sized vector as input and produce a fixed-sized vector as output
- fixed amount of computational steps
- recurrent nets allow us to operate over *sequences* of vectors

Use cases

- Video: sequence understanding
- Audio: speech transcription
- Text: natural language processing

Recurrent neuron

- x_t : Input at time t
- h_{t-1} : State at time $t-1$



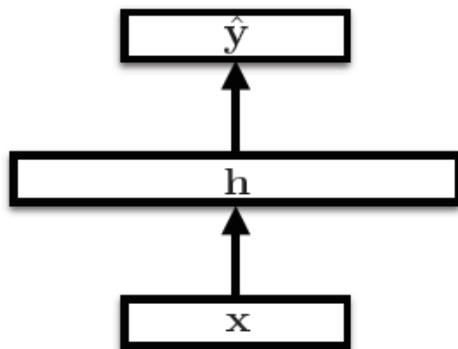
$$h_t = f(W_h h_{t-1} + W_x x_t)$$

Recurrent Neural Networks

Feed-forward NN

$$\mathbf{h} = g(\mathbf{V}\mathbf{x} + \mathbf{c})$$

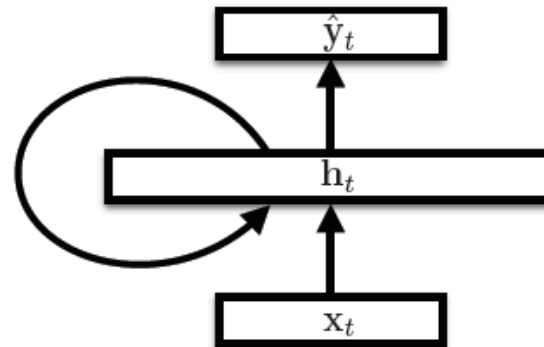
$$\hat{\mathbf{y}} = \mathbf{W}\mathbf{h} + \mathbf{b}$$



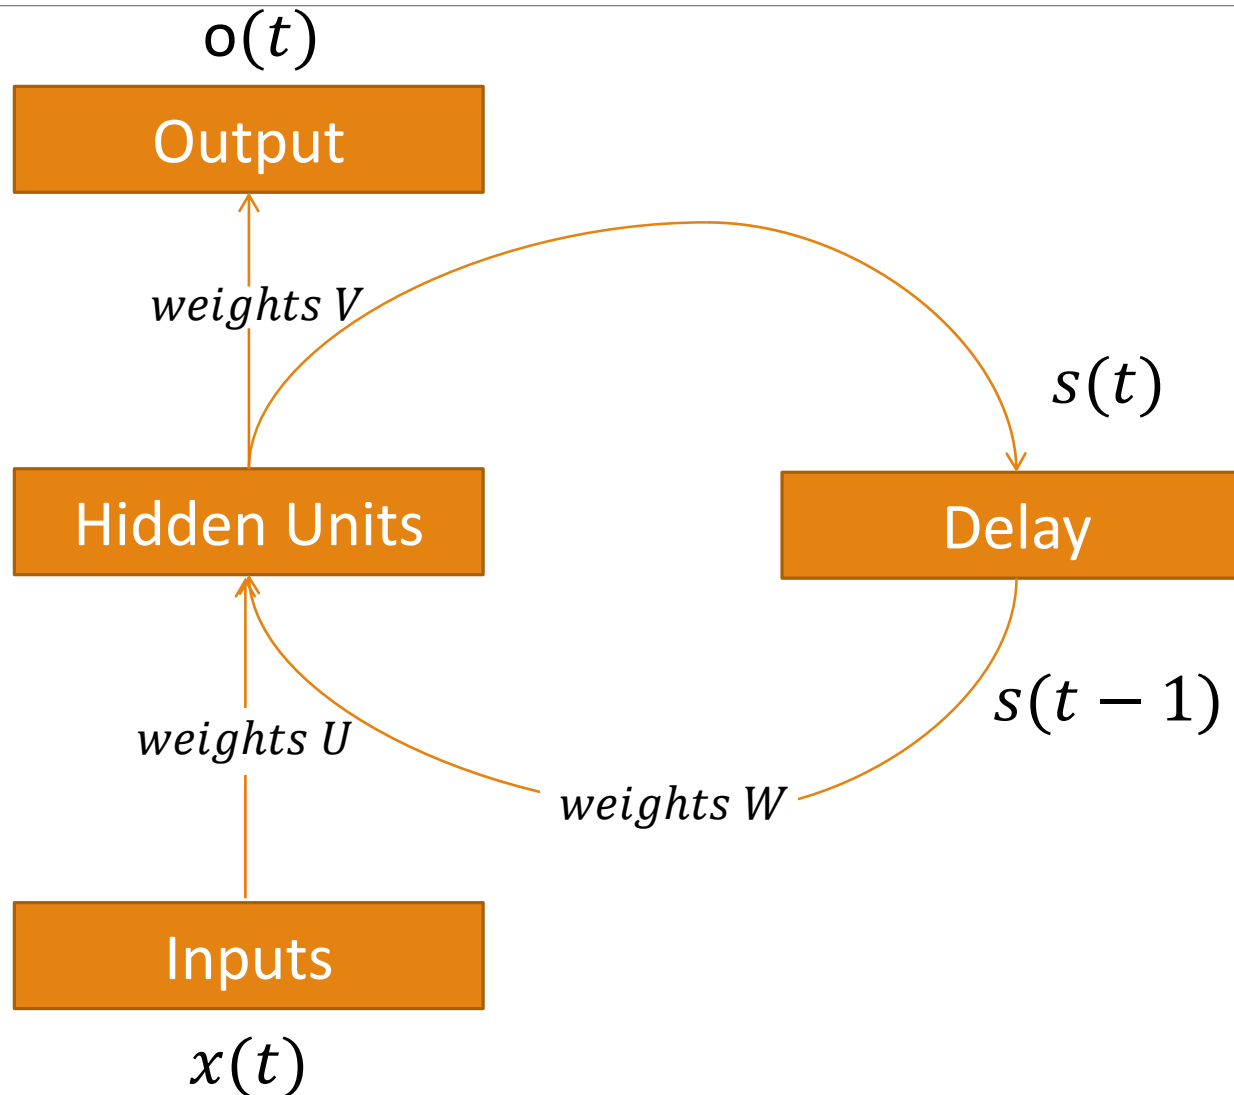
Recurrent NN

$$\mathbf{h}_t = g(\mathbf{V}\mathbf{x}_t + \mathbf{U}\mathbf{h}_{t-1} + \mathbf{c})$$

$$\hat{\mathbf{y}}_t = \mathbf{W}\mathbf{h}_t + \mathbf{b}$$

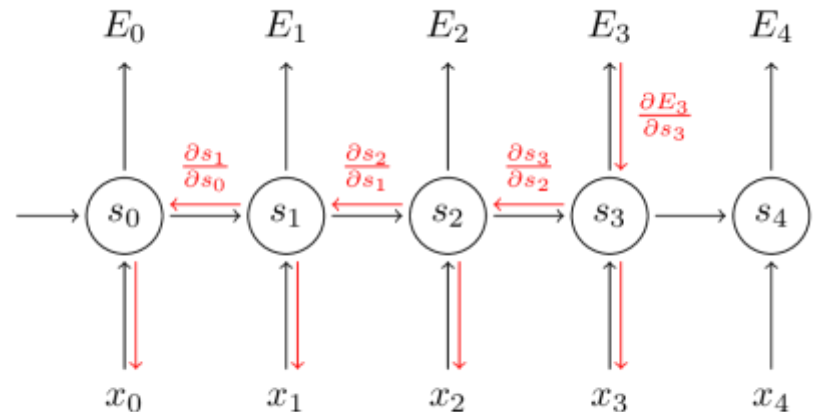
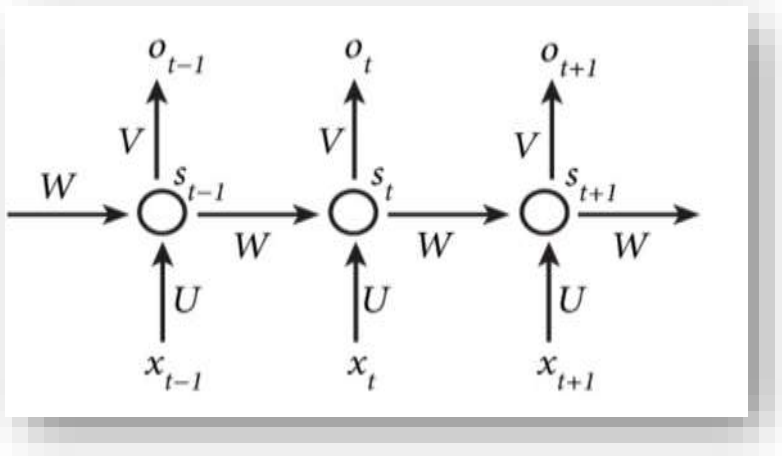


RNN Architecture



Unfolding RNNs

- Each node represents a layer of network units at a single time step.
- The same weights are reused at every time step.



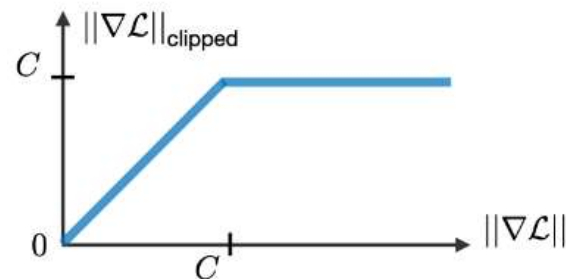
Training RNNs

Loss function: $\mathcal{L}(\hat{y}, y) = \sum_{t=1}^{T_y} \mathcal{L}(\hat{y}^{<t>}, y^{<t>})$

Backpropagation through time $\frac{\partial \mathcal{L}^{(T)}}{\partial W} = \sum_{t=1}^T \frac{\partial \mathcal{L}^{(T)}}{\partial W} \Big|_{(t)}$

Vanishing/exploding gradient: failure to capture long-range dependencies

Gradient clipping



RNNs pros and cons

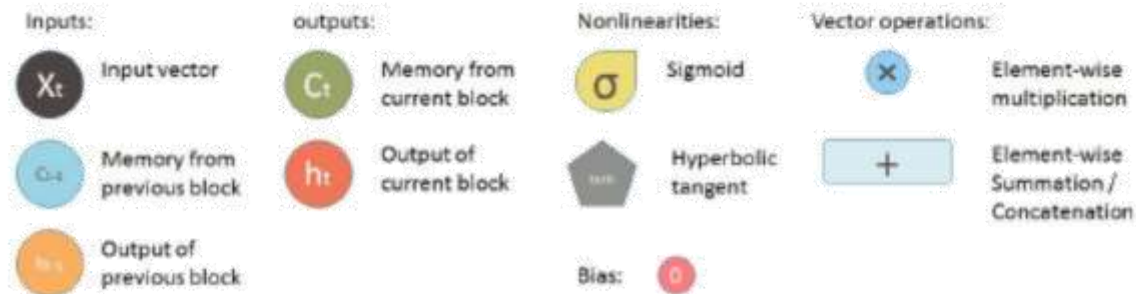
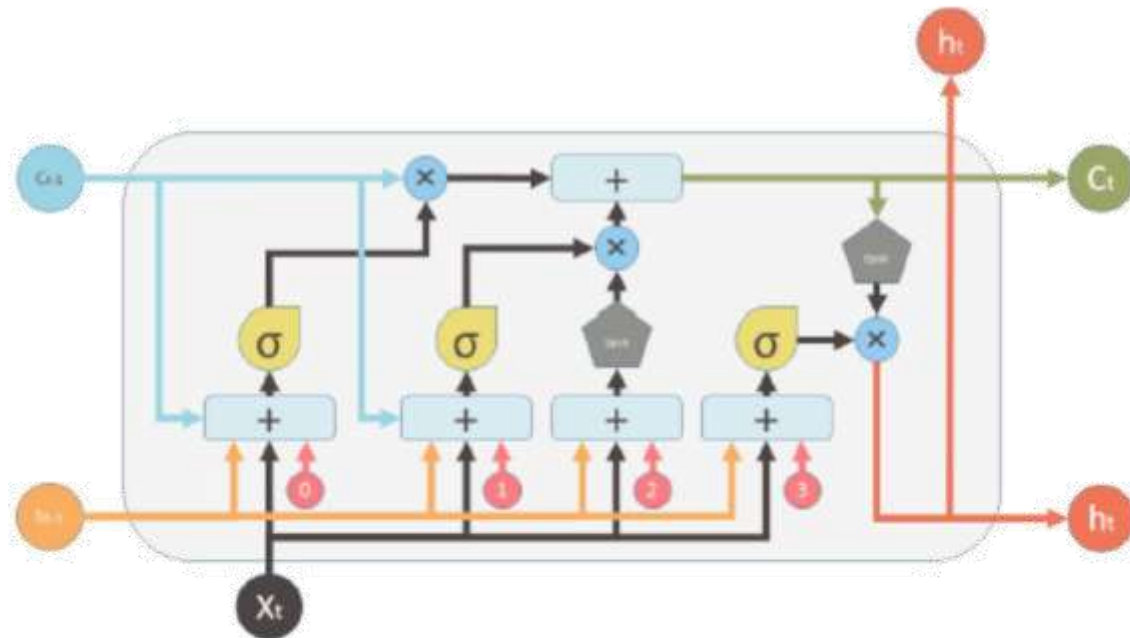
ADVANTAGES

- Possibility of processing input of any length
- Model size not increasing with size of input
- Computation takes into account historical information
- Weights are shared across time

DISADVANTAGES

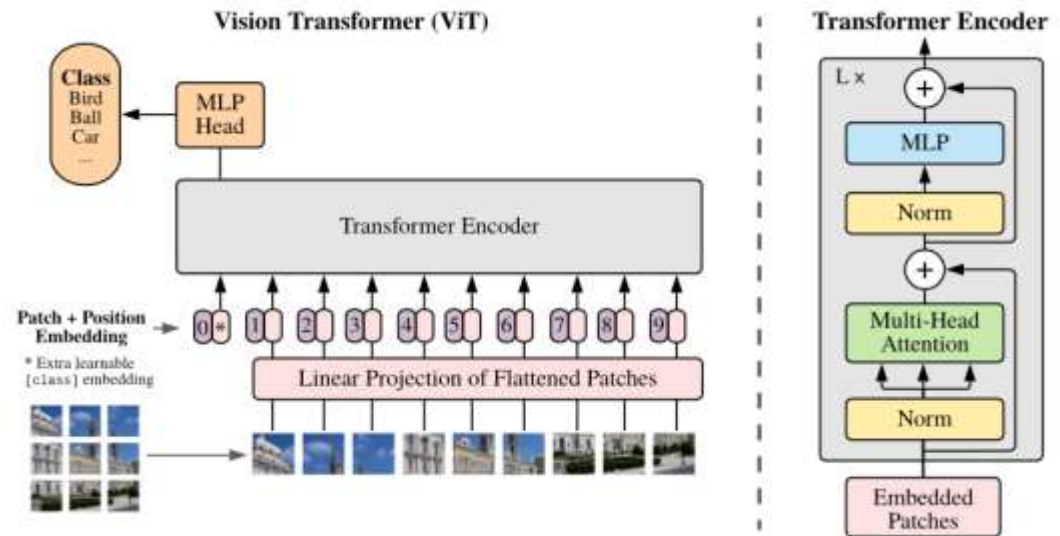
- Computation being slow
- Difficulty of accessing information from a long time ago
- Cannot consider any future input for the current state

Long Short-Term Memory Nets (LSTMs)



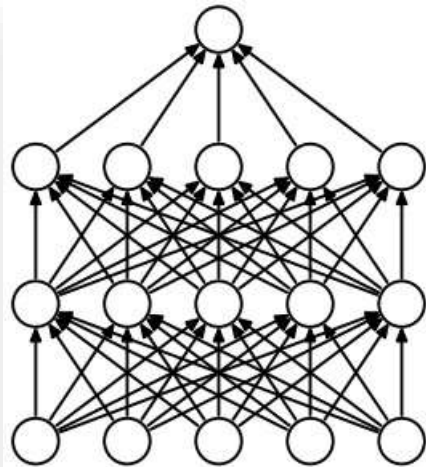
Vision Transformer (ViT)

- Split an image into patches (fixed sizes)
- Flatten the image patches
- Create lower-dimensional linear embeddings from these flattened image patches
- Include positional embeddings
- Feed the sequence as an input to a state-of-the-art transformer encoder
- Pre-train the ViT model with image labels, which is then fully supervised on a big dataset
- Fine-tune the downstream dataset for image classification

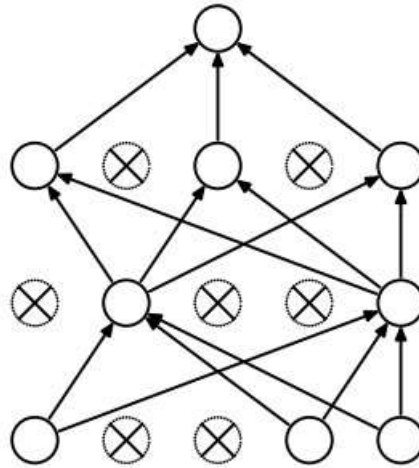


Dosovitskiy, Alexey, et al. "An image is worth 16x16 words: Transformers for image recognition at scale." *arXiv preprint arXiv:2010.11929* (2020).

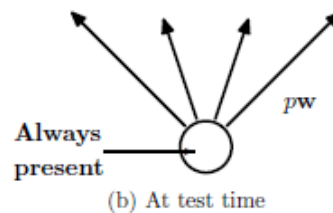
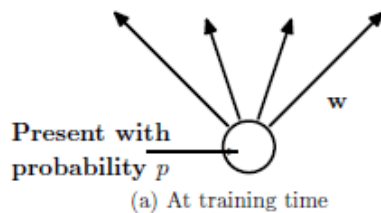
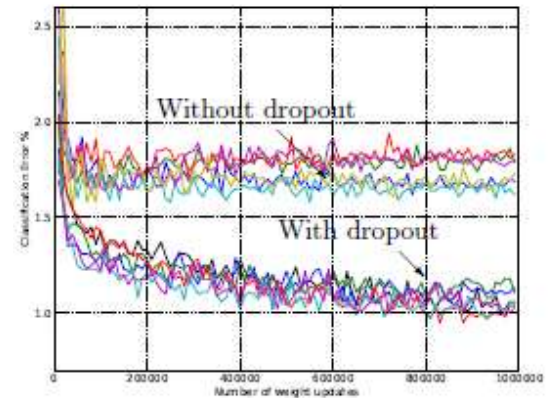
Dropout



(a) Standard Neural Net



(b) After applying dropout.



Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting." *Journal of machine learning research* 15.1 (2014): 1929-1958.

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

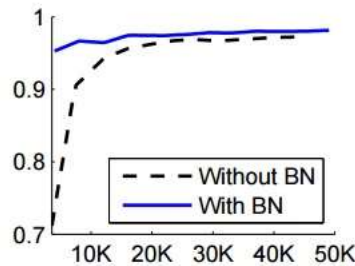
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

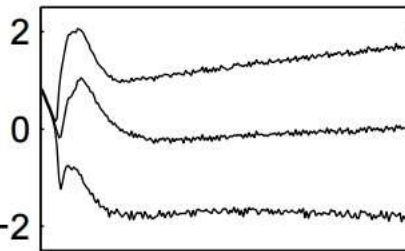
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

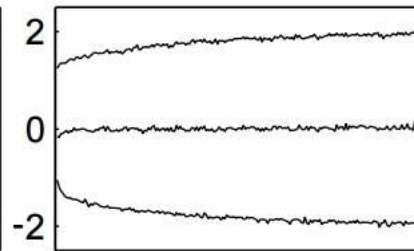
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$



(a)

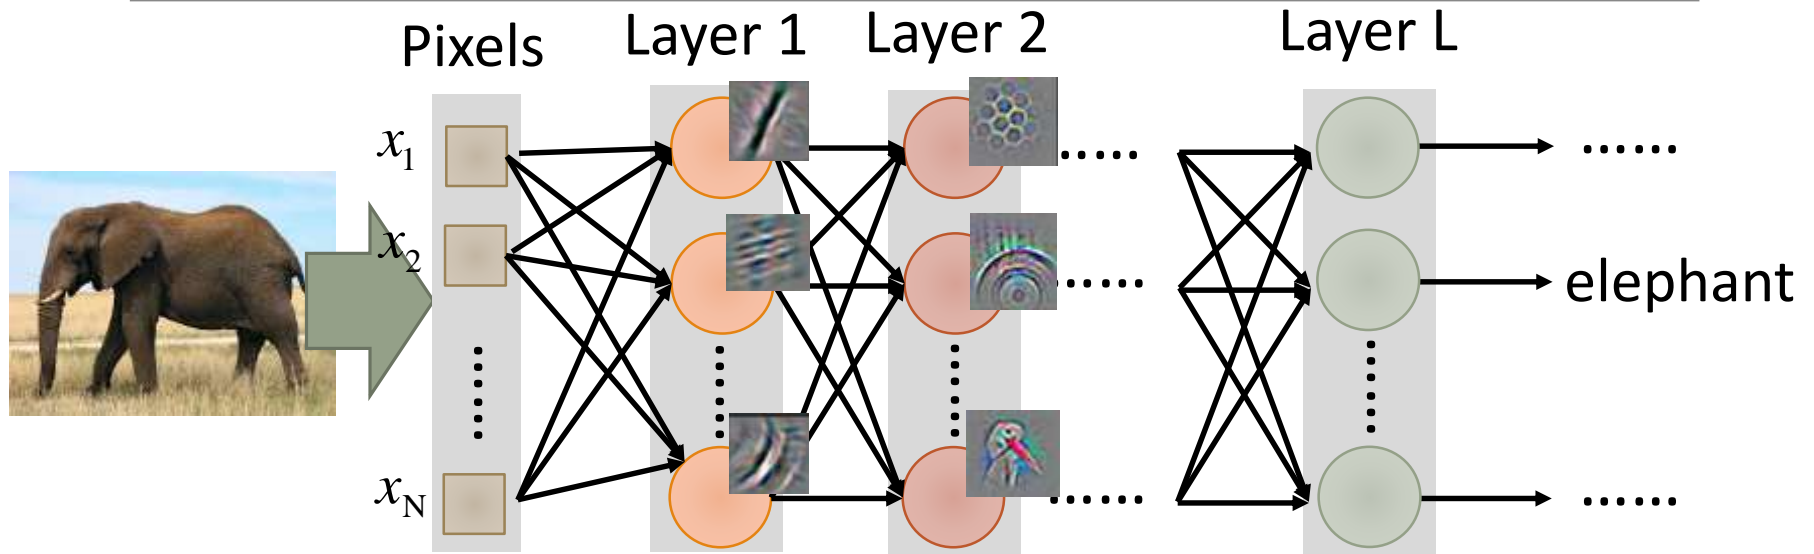


(b) Without BN

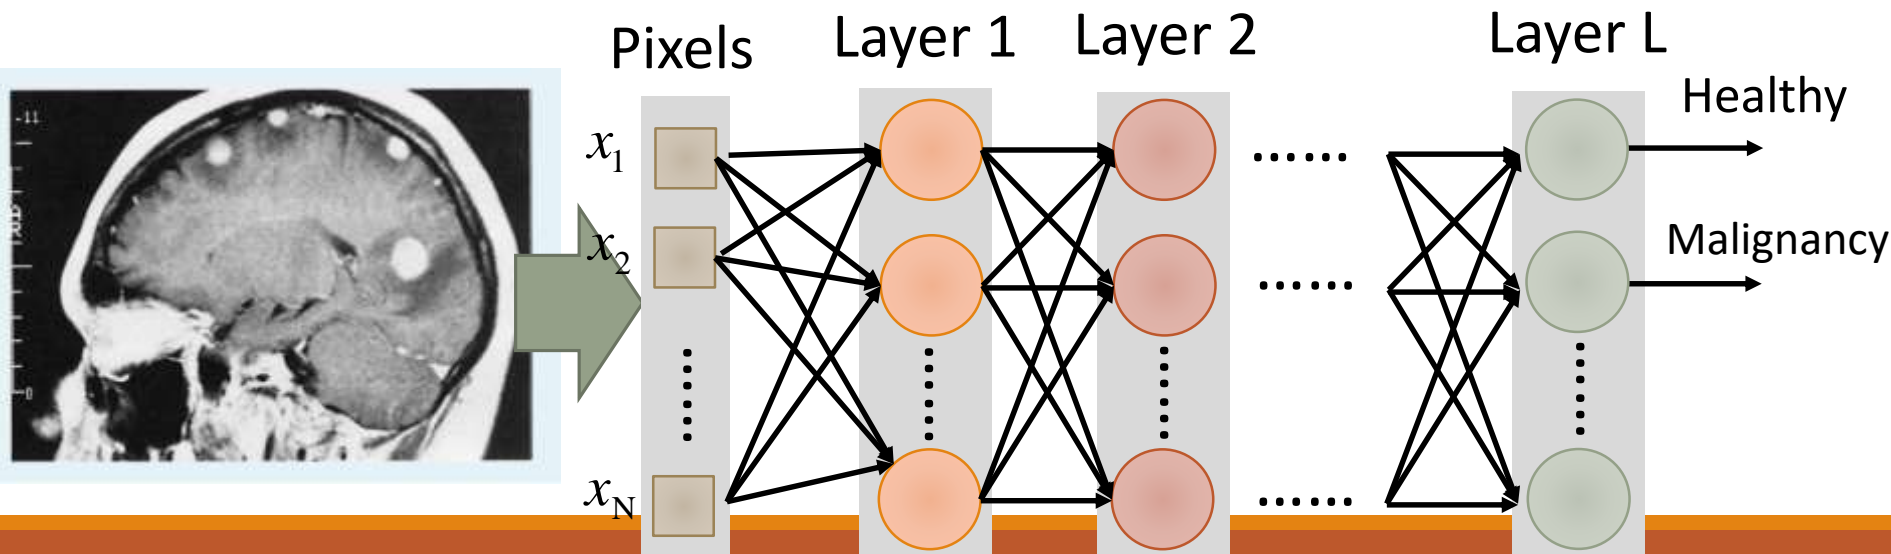
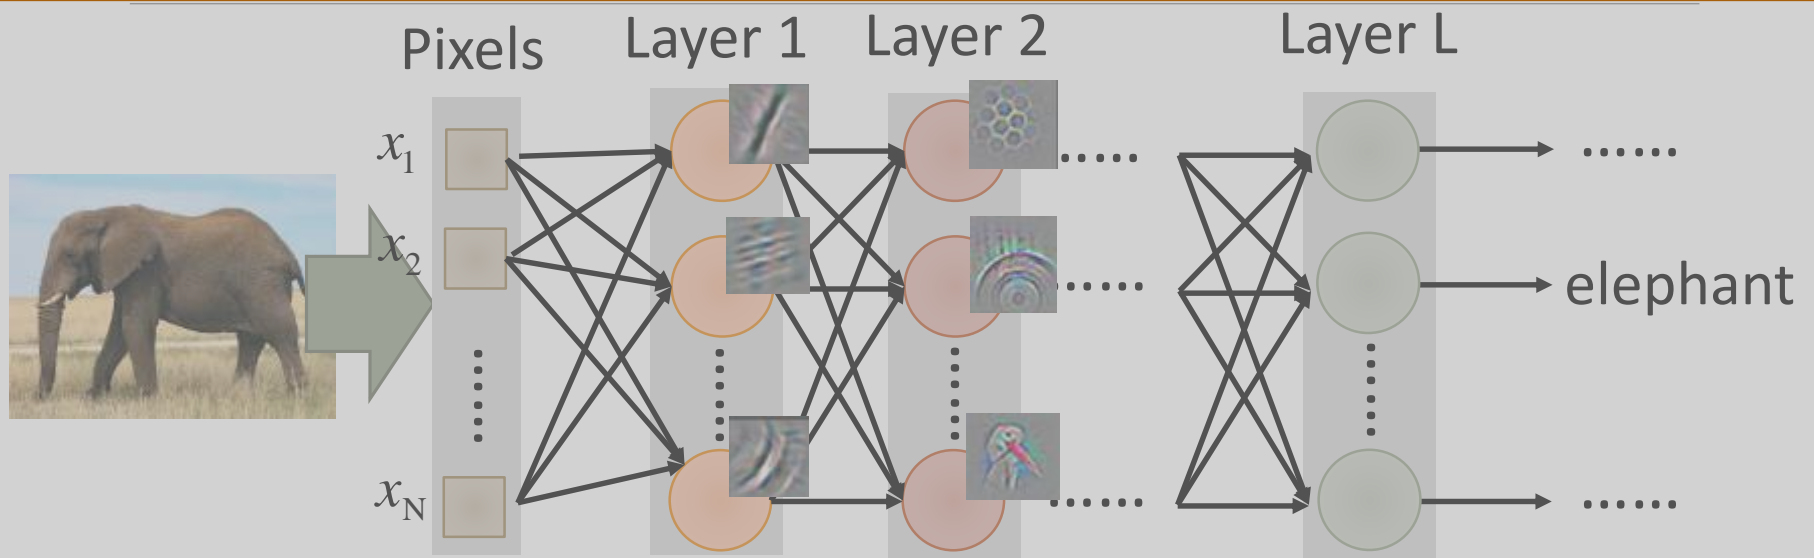


(c) With BN

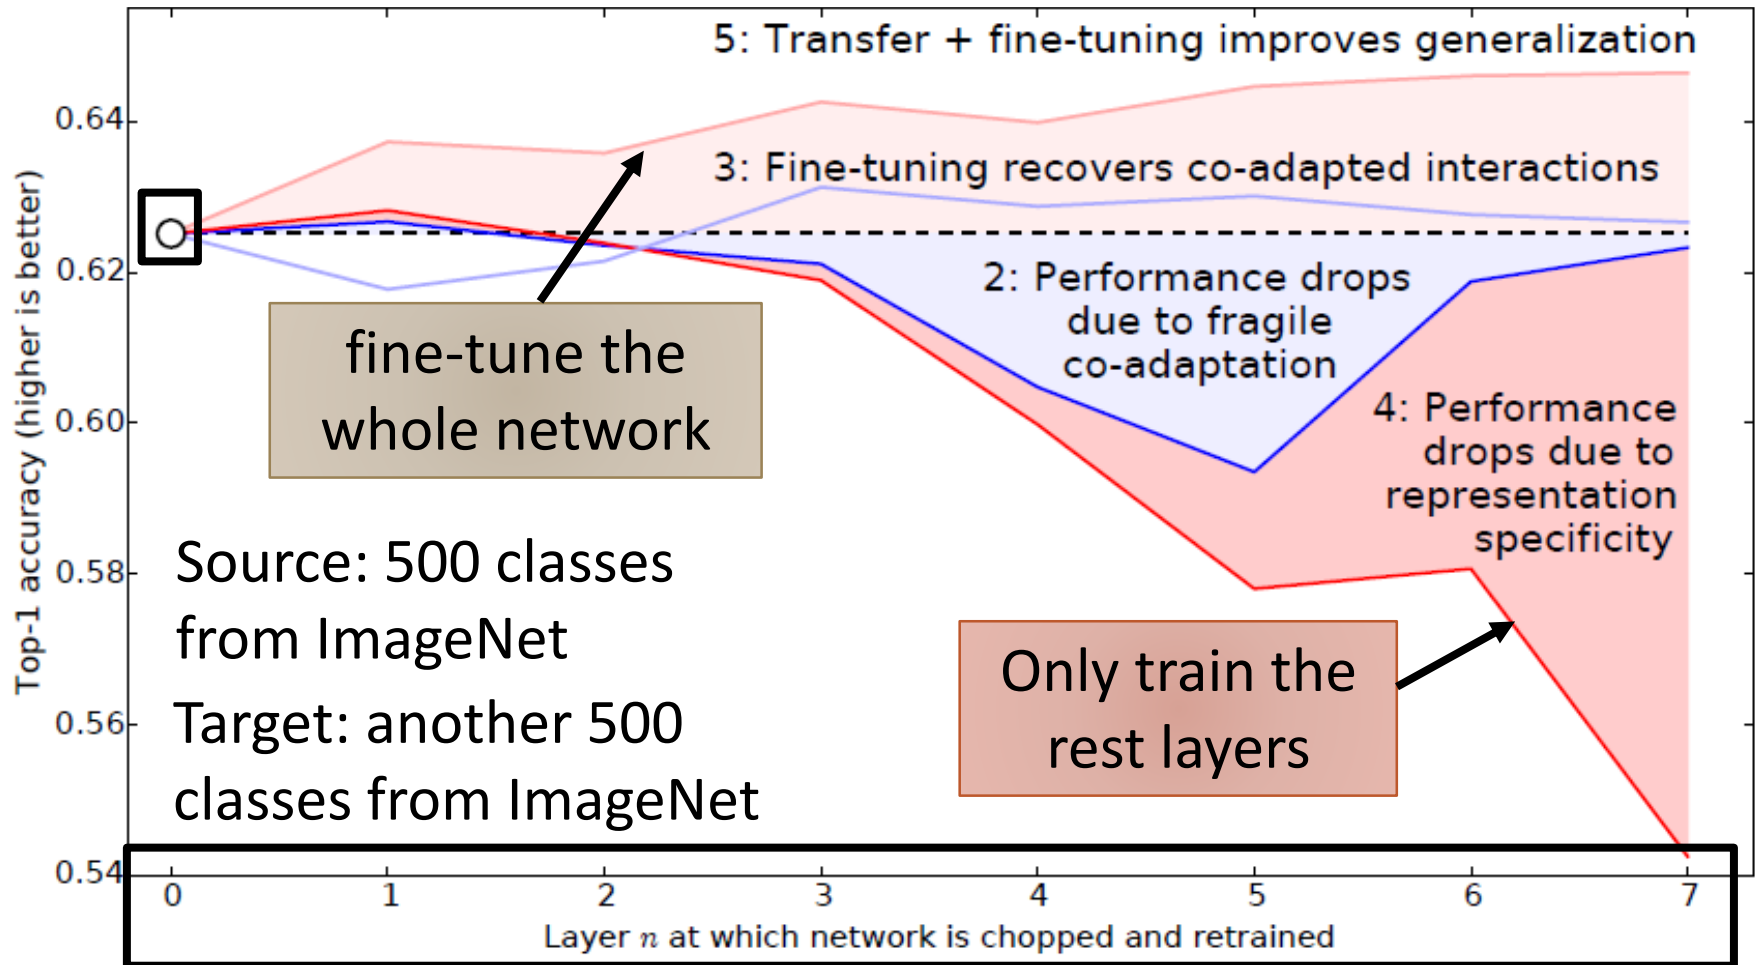
Transfer Learning



Transfer Learning

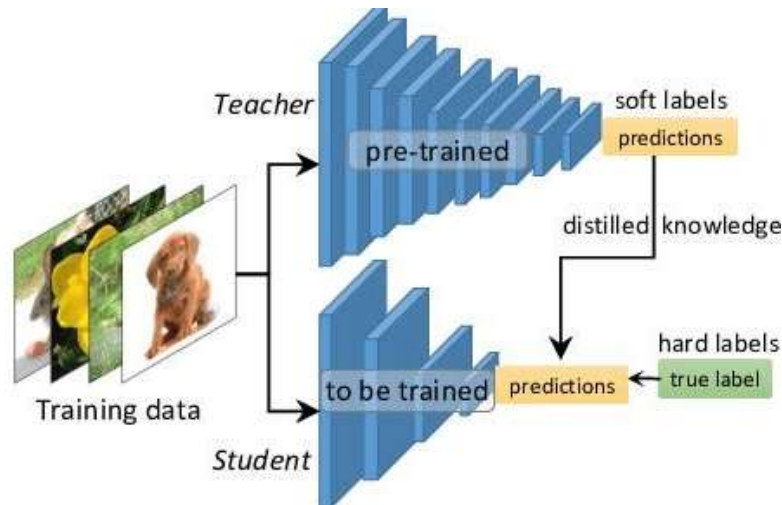
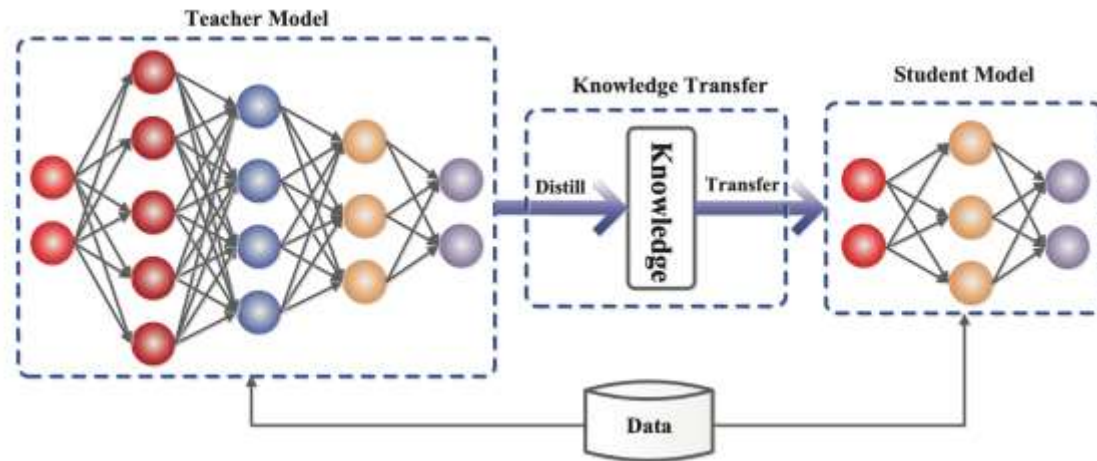


Layer Transfer - Image



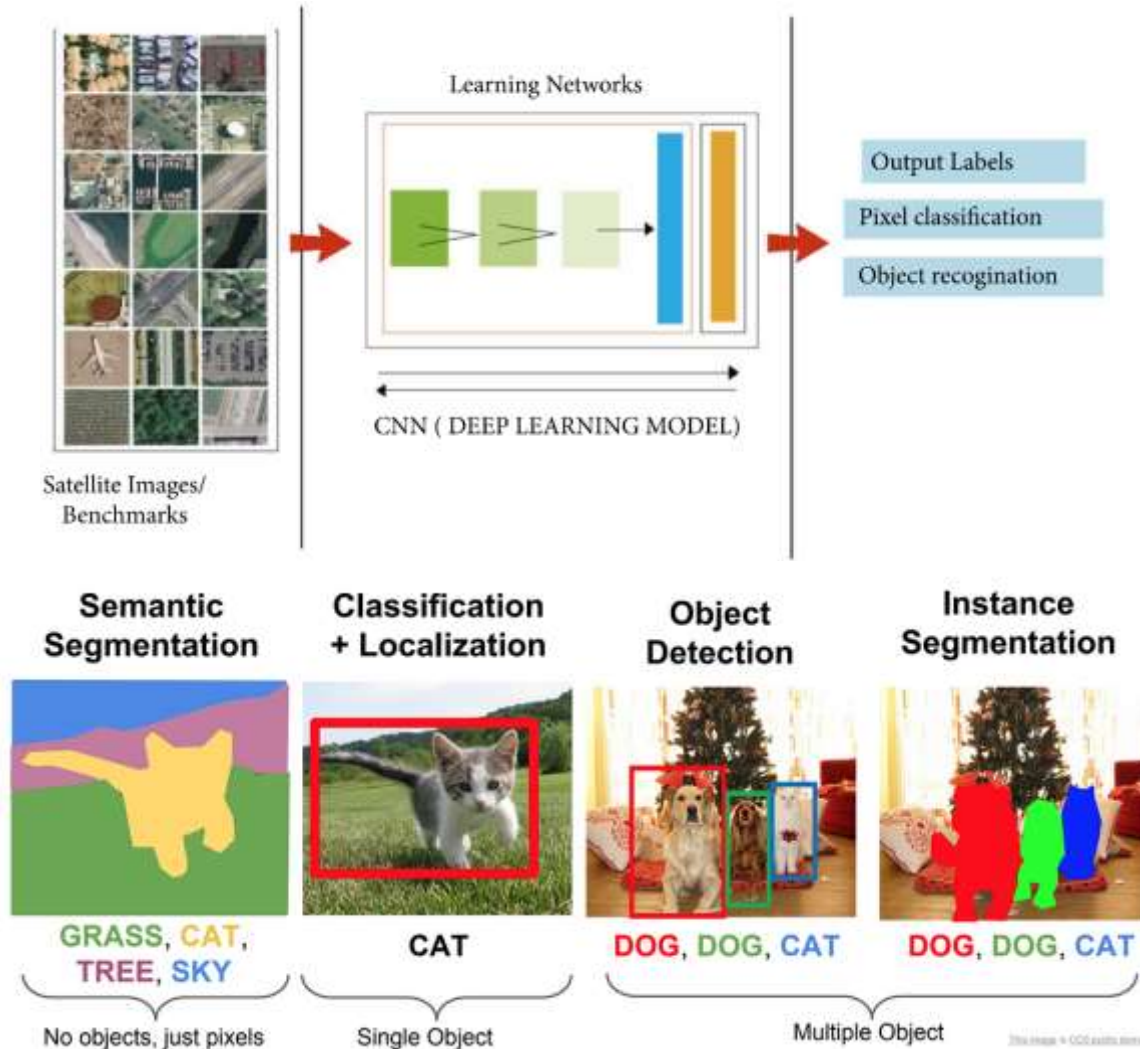
J. Yosinski, J. Clune, Y. Bengio, H. Lipson, "How transferable are features in deep neural networks?", NIPS, 2014

Knowledge distillation



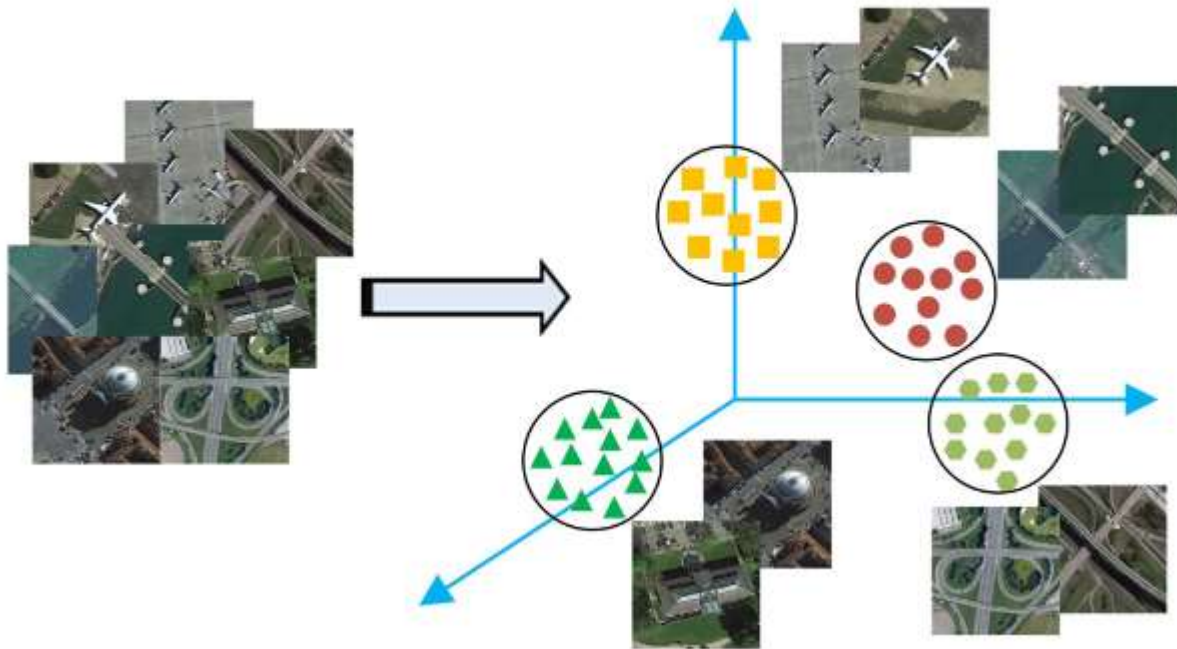
<https://neptune.ai/blog/knowledge-distillation>

Applications in RS



Multi-class classification

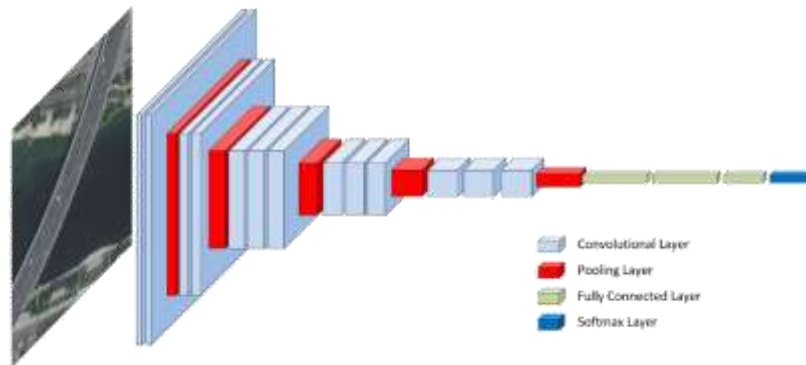
Scene classification



Remote Sensing Image Scene
Classification Meets Deep Learning: Challenges,
Methods, Benchmarks, and Opportunities

Performance on UCMerced





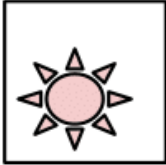

CNN-based	GBRCN [102]	2015	IEEE TGRS	-	94.53
	LPCNN [103]	2016	JARS	-	89.90
	Fusion by Addition [109]	2017	IEEE TGRS	-	97.42±1.79
	ARCNet-VGG16 [74]	2018	IEEE TGRS	96.81±0.14	99.12±0.40
	MSCP [112]	2018	IEEE TGRS	-	98.36±0.58
	D-CNNs [73]	2018	IEEE TGRS	-	98.93±0.10
	MCNN [116]	2018	IEEE TGRS	-	96.66±0.9
	ADSSM [138]	2018	IEEE TGRS	-	99.76±0.24
	FACNN [113]	2019	IEEE TGRS	-	98.81±0.24
	SF-CNN [118]	2019	IEEE TGRS	-	99.05±0.27
	SCCov [123]	2019	IEEE TNNLS	-	99.05±0.25
	RSFJR [117]	2019	IEEE TGRS	97.21±0.65	-
	GBN [119]	2019	IEEE TGRS	97.05±0.19	98.57±0.48
	ADFF [139]	2019	Remote Sensing	96.05±0.56	97.53±0.63
	CNN-CapsNet [140]	2019	Remote Sensing	97.59±0.16	99.05±0.24
	Siamese ResNet50 [141]	2019	IEEE GRSL	90.95	94.29



Remote Sensing Image Scene Classification Meets Deep Learning: Challenges, Methods ,Benchmarks, and Opportunities

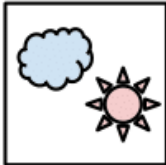


Multi-class vs. Multi-label

Multi-Class

C = 3	Samples
  	 [0 0 1]  [1 0 0]  [0 1 0]

Classes are mutually exclusive

Multi-Label

Samples
 [1 0 1]  [0 1 0]  [1 1 1]

No restrictions on the number of associated labels per sample

Multiclass vs Multi-label

Land Cover, Scene Classification

Dense Residential



Airplane



Sparse Residential



Forest



Scene characterization

**Buildings, Cars,
Pavement, Trees**



**Airplane, Cars,
Grass, Pavement**



**Buildings, Cars, Chaparral,
Pavement, Sand, Trees**



Trees



A CNN based approach

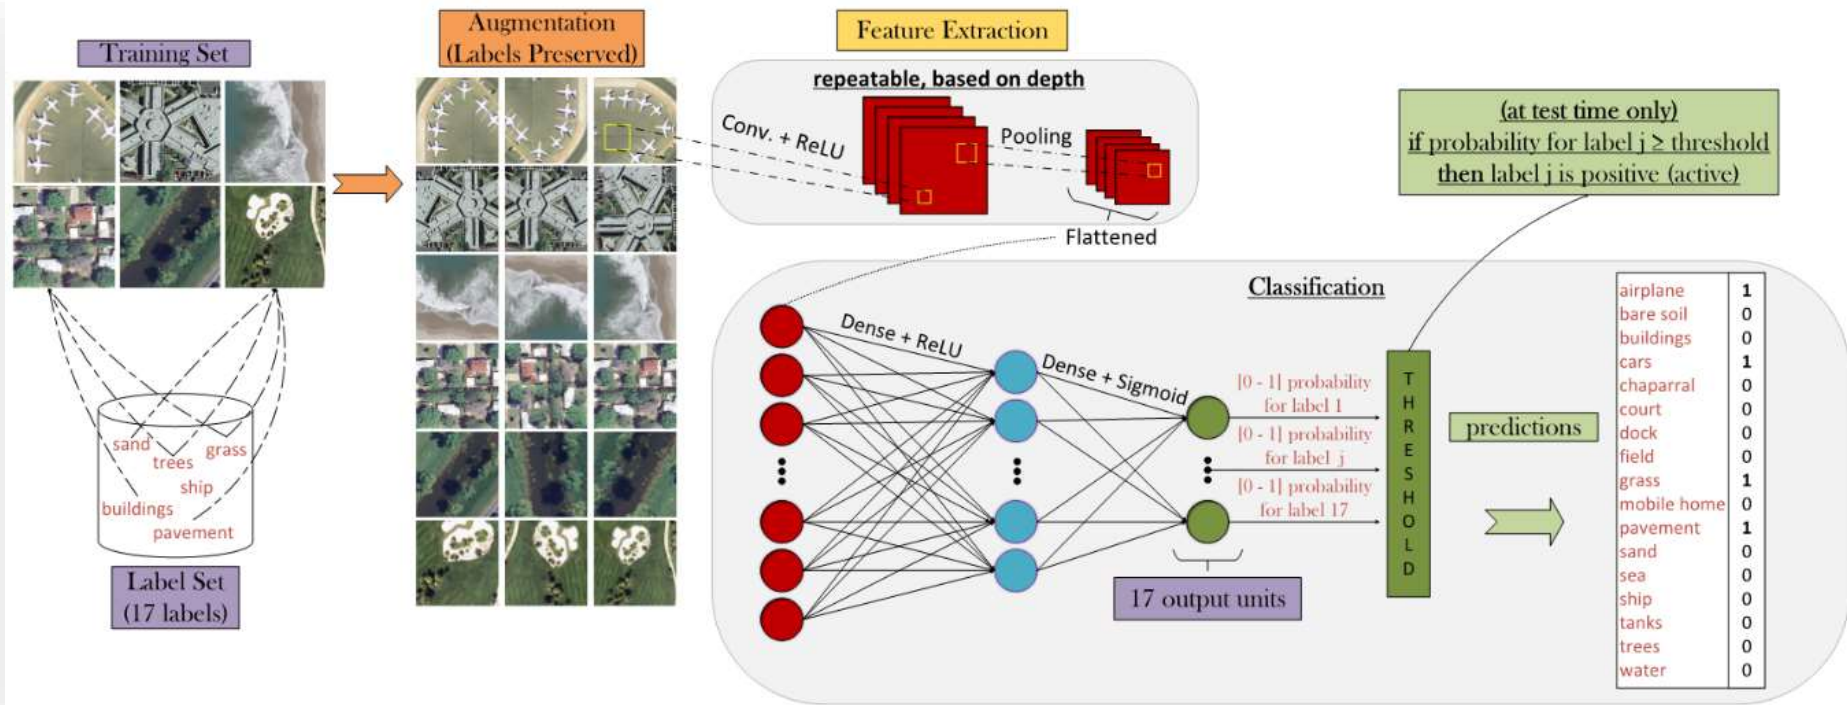
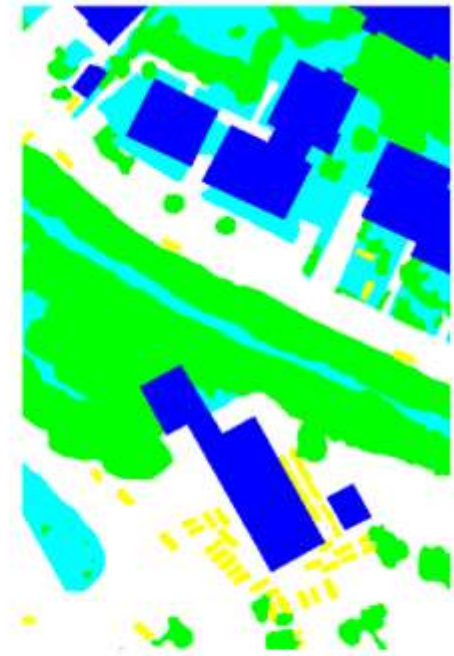
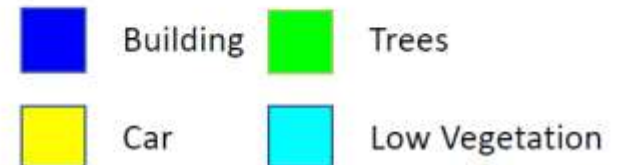
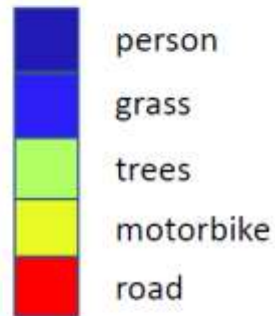


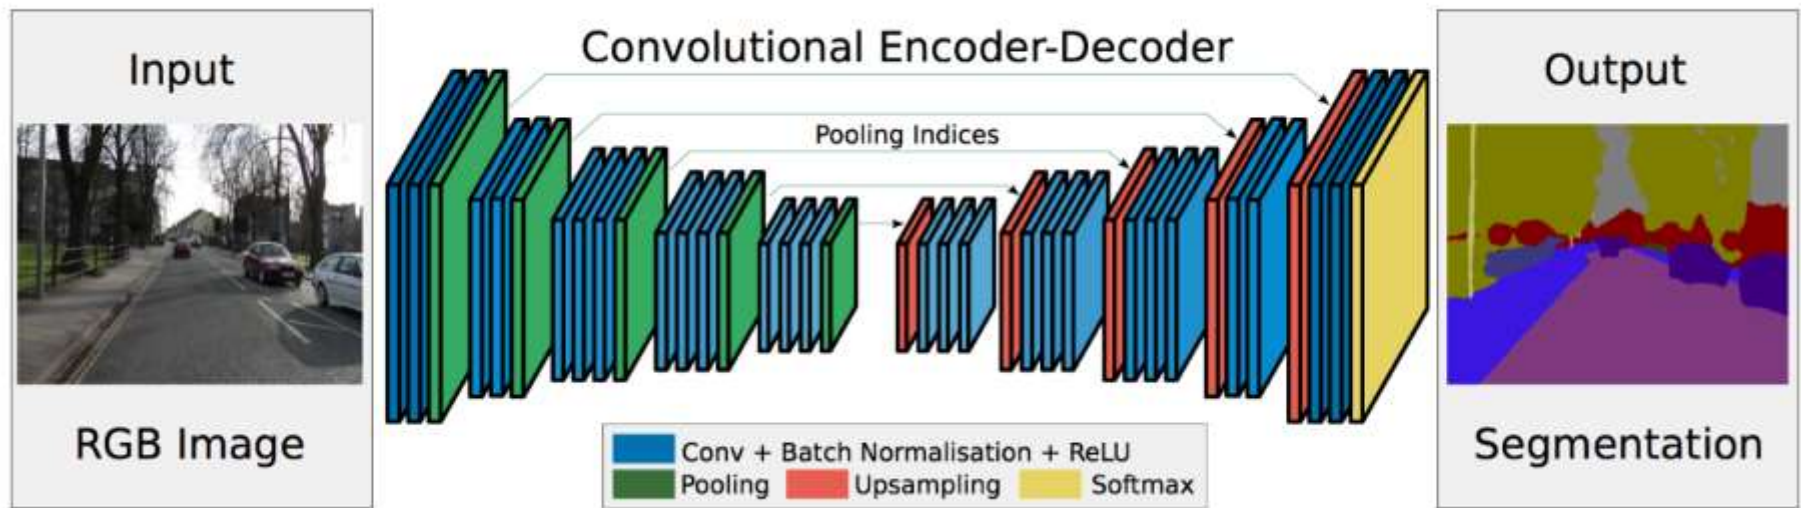
Image segmentation



SegNet

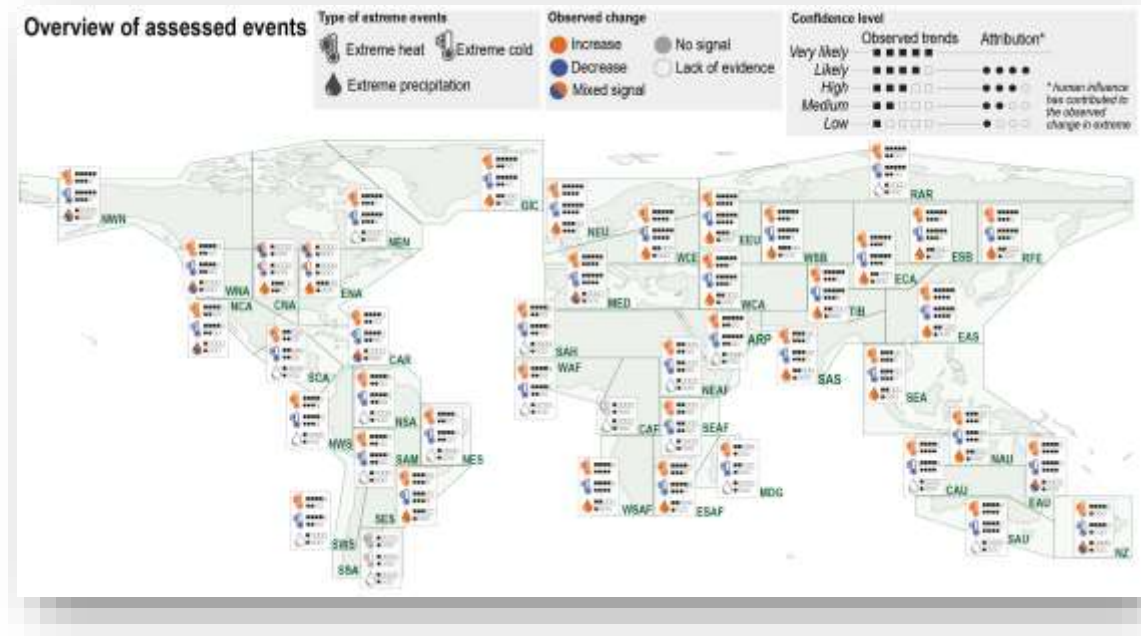
Fully convolutional networks (FCN) à no dense layers

Can handle different input sizes

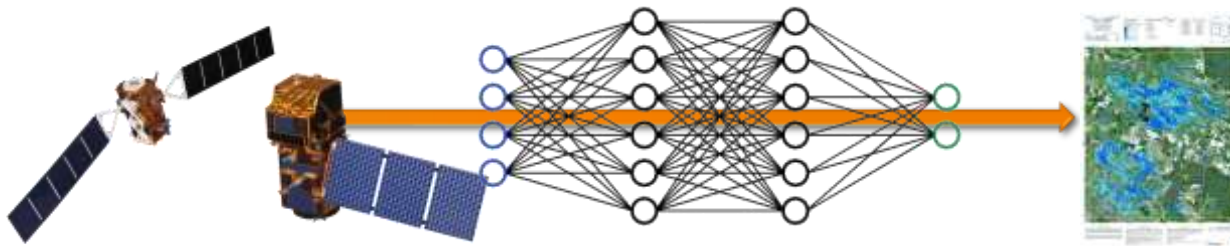


Badrinarayanan, Vijay, Alex Kendall, and Roberto Cipolla. "Segnet: A deep convolutional encoder-decoder architecture for image segmentation." IEEE Transactions on pattern analysis and machine intelligence 39.12 (2017): 2481-2495.

Extreme weather events



[IPCC Sixth Assessment Report](#)



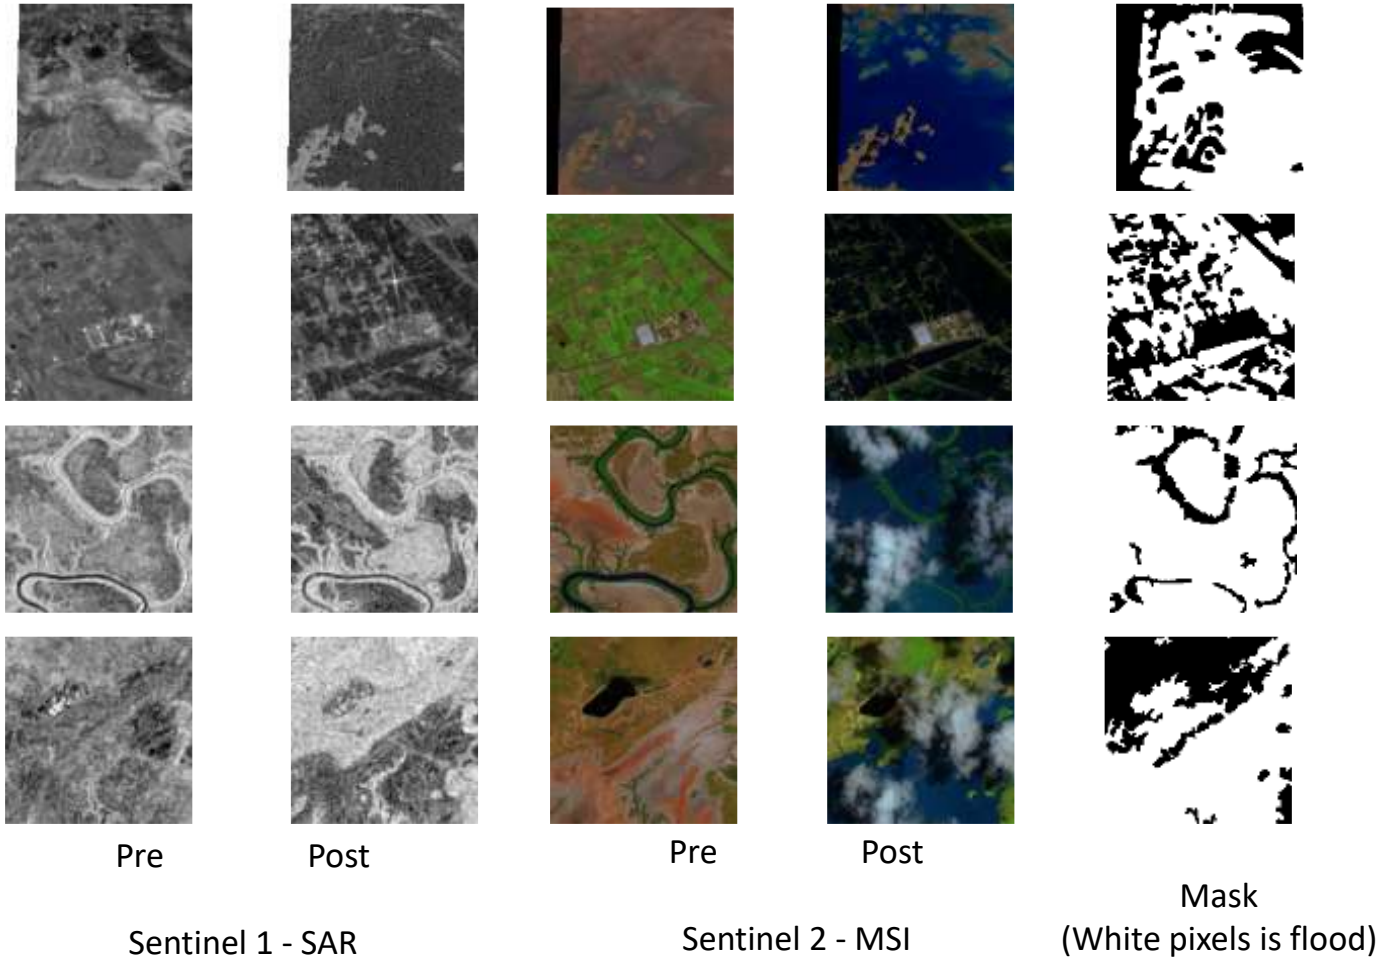
OMBRIA dataset



EMS ID	Country	Date 1	Date 2	UTM Zone
271	Greece	01/05/2017	28/02/2018	34 N
273	Albania	01/05/2017	11/03/2018	34 N
275	Croatia	01/05/2017	22/03/2018	33 N
279	Spain	01/05/2017	15/04/2018	30 N
324	France	01/05/2018	16/10/2018	31 N
342	Australia	15/04/2018	13/02/2019	54 S
388	Spain	01/05/2019	14/09/2019	30 N
416	France	01/05/2019	15/12/2019	30 N
417	Portugal	01/05/2019	23/12/2019	29 N
419	Iran	01/05/2019	13/01/2020	41 N
422	Spain	01/05/2019	26/01/2020	31 N
424	Madagascar	01/05/2019	29/01/2020	39 S
429	Ireland	01/05/2019	23/02/2020	29 N
441	Finland	01/05/2019	04/06/2020	34 N
465	Greece	01/05/2020	20/09/2020	31 N
466	Niger	01/05/2020	27/09/2020	32 N
468	Italy	01/05/2020	10/10/2020	32 N
470	Togo	01/05/2020	17/10/2020	31 N
482	Honduras	01/05/2020	22/11/2020	17 N
492	France	01/05/2020	02/01/2021	30 N
501	Albania	01/05/2020	15/02/2021	35 N
507	Timor	01/05/2020	06/04/2021	51 S
514	Guyana	01/05/2020	06/06/2021	21 S

Drakonakis, G. I., Tsagkatakis, G., Fotiadou, K., & Tsakalides, P. "Ombrianet—supervised flood mapping via convolutional neural networks using multitemporal sentinel-1 and sentinel-2 data fusion". *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 2022.

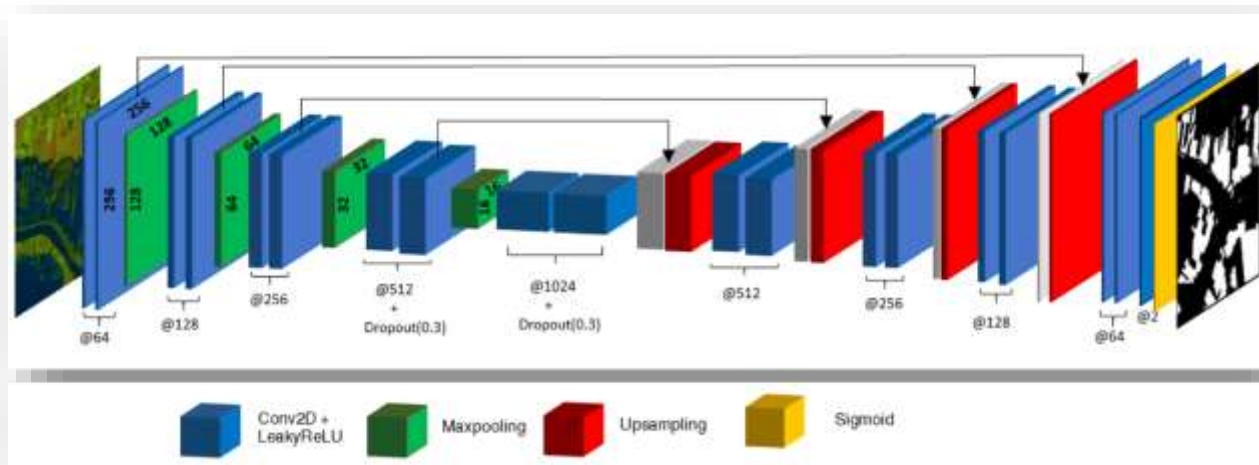
OMBRIA dataset



2D UNET

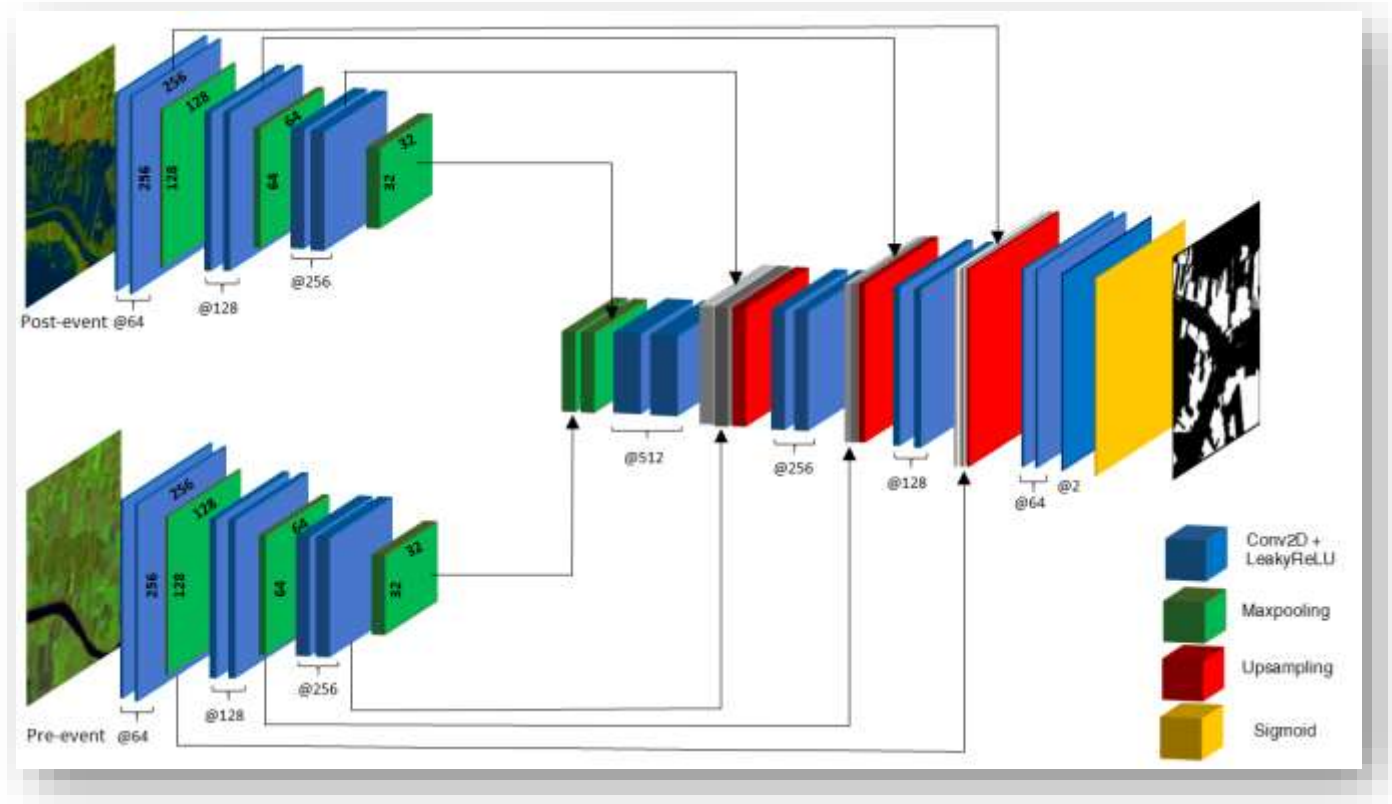
U-NET:

- Image -> image
- Image segmentation

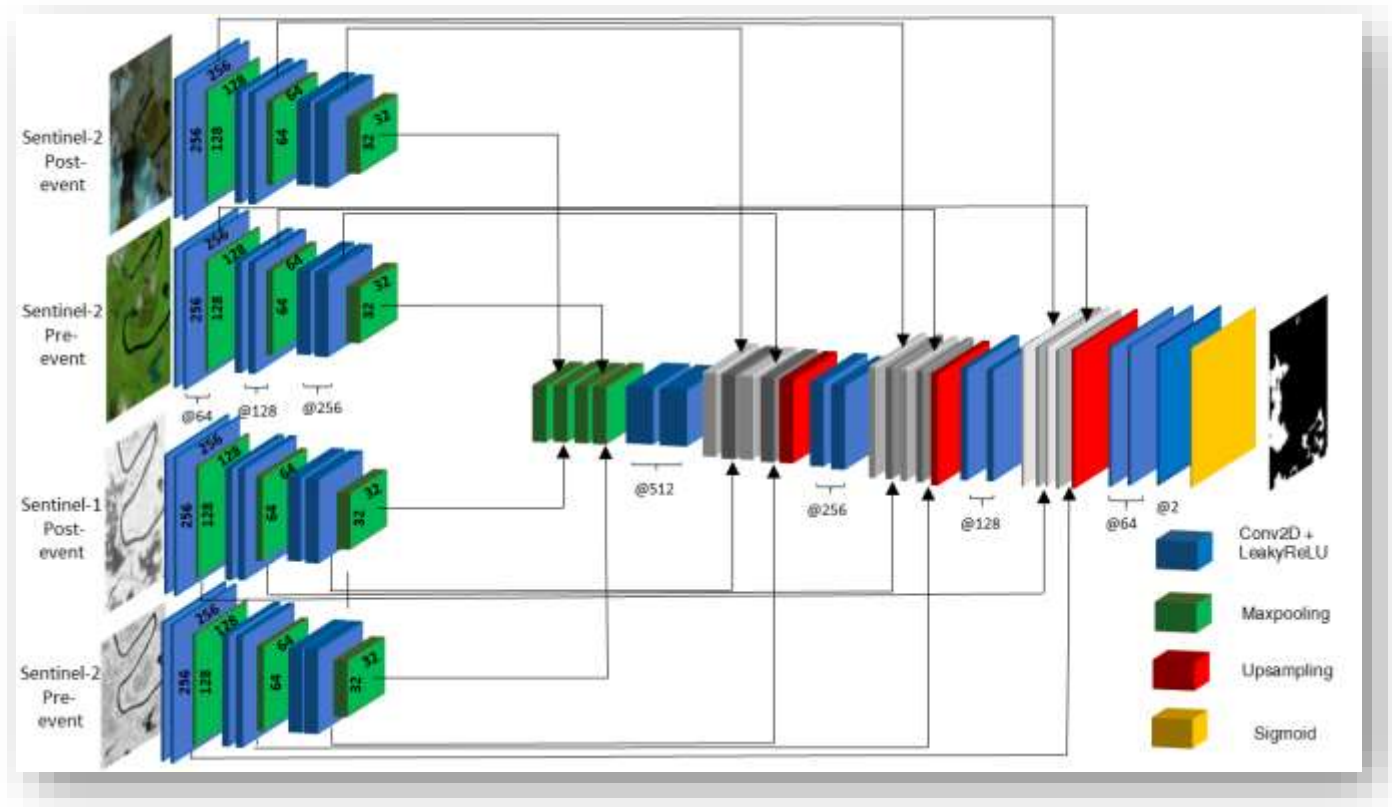


Ronneberger, O., Fischer, P., & Brox, T. (2015). U-net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2015*.

3D UNET



4D UNET



OMBRIA NET

Methods	PA	IoU	FW IoU
Otsu's Thresholding (Sentinel-2)	0.7930	0.333	0.6889
Multimodal SVM	0.8504	0.6245	0.7631
U-Net (Sentinel-1)	0.7925	0.5734	0.6971
U-Net (Sentinel-2)	0.8251	0.5418	0.7221
Bitemporal OmbriaNet (Sentinel-1)	0.7203	0.5181	0.6229
Bitemporal OmbriaNet (Sentinel-2)	0.8733	0.6457	0.7919
Multimodal OmbriaNet (Sentinel-1 & Sentinel-2)	0.9010	0.7236	0.8330



(a) S-1 (pre)



(b) S-1 (post)



(c) S-2 (pre)



(d) S-2 (post)



(e) S-1 (63.8 %)



(f) S-2
(78.3%)

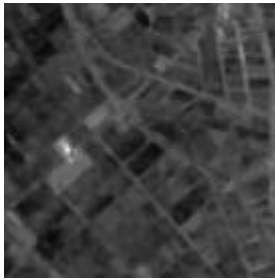


(g) OmbriaNet (89.1%)

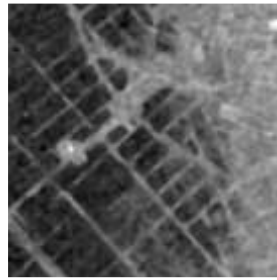


(h) Ground Truth
(White pixels is flood)

Performance



(a) S-1 (Pre-event)



(b) S-1 (Post-event)



(c) S-2 (Pre-event)



(d) S-2 (Post-event)



(e) S-1 U-Net
(22.19%)



(f) S-2 U-Net
(59.13%)



(g) OmbriaNet
(81.90%)



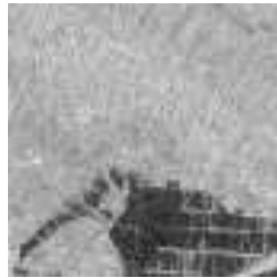
(h) Ground Truth

Performance

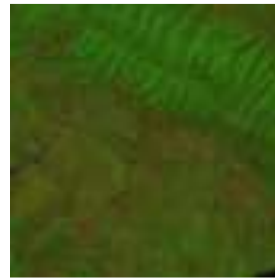
Comparison of selected sample from ID501 flood in Albania (IoU metric score)



(a) S-1 (Pre-event)



(b) S-1 (Post-event)



(c) S-2 (Pre-event)



(d) S-2 (Post-event)



(e) Sentinel-1 U-Net
(63.82 %)



(f) Sentinel-2 U-Net
(78:30%)



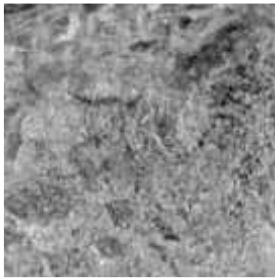
(g) OmbriaNet
(89:14%)



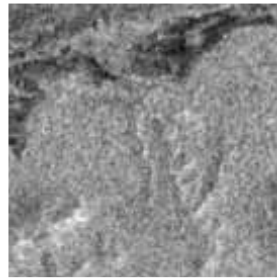
(h) Ground Truth (White
pixels is flood)

Performance

Comparison of selected sample from ID507 flood in Timor



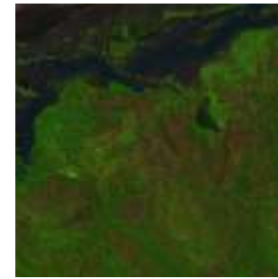
(a) S-1 (Pre-event)



(b) S-1 (Post-event)



(c) S-2 (Pre-event)



(d) S-2 (Post-event)



(e) Sentinel-1
U-Net (69:50 %)



(f) Sentinel-2 U-
Net (79:08%)



(g) OmbriaNet
(79:44%)

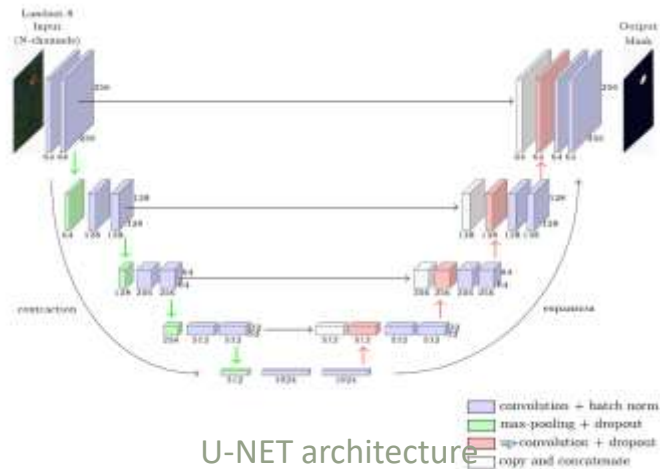


(h) Ground Truth
(White pixels is
flood)

Active fire detection

$((R_{75} > 2.5) \text{ and } (\rho_7 - \rho_5 > 0.3) \text{ and } (\rho_7 > 0.5)) \text{ or } ((\rho_6 > 0.8) \text{ and } (\rho_1 < 0.2) \text{ and } (\rho_5 > 0.4 \text{ or } \rho_7 < 0.1))$

Schroeder et al. (2016)



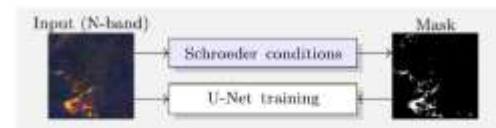
U-NET architecture

$(R_{76} \geq 1.4) \text{ and } (R_{75} \geq 1.4) \text{ and } (\rho_7 \geq 0.15)$

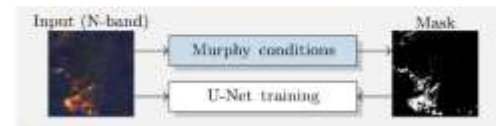
Murphy et al. (2016)

$\rho_4 \leq 0.53 \rho_7 - 0.214$

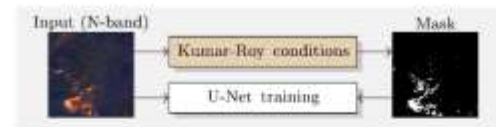
Kumar-Roy



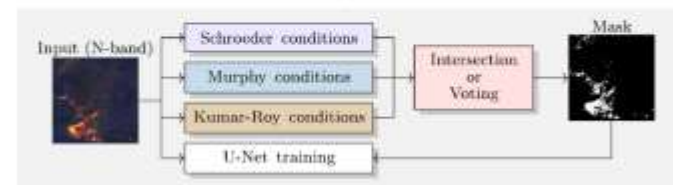
(a)



(b)



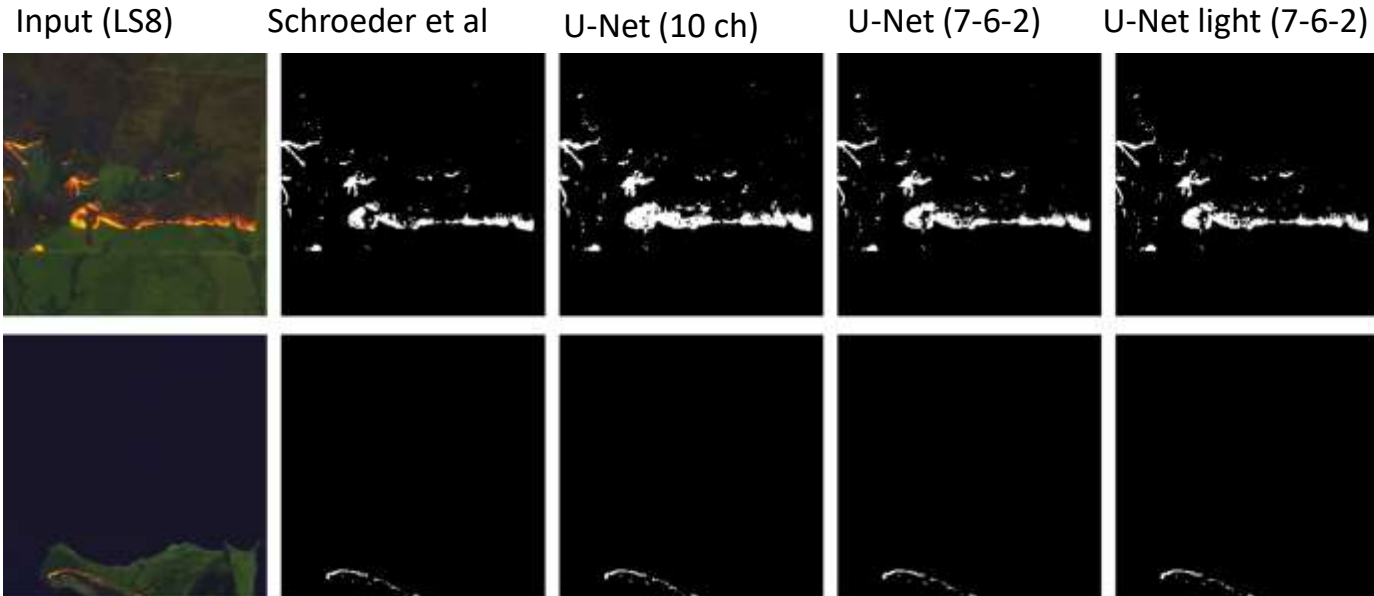
(c)



(d)

de Almeida Pereira, Gabriel Henrique, et al. "Active fire detection in Landsat-8 imagery: A large-scale dataset and a deep-learning study." *ISPRS Journal of Photogrammetry and Remote Sensing* 178 (2021): 171-186.

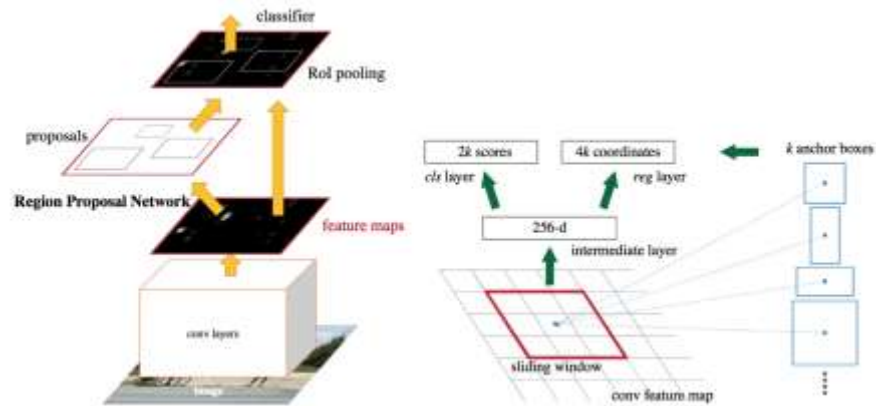
Active fire detection



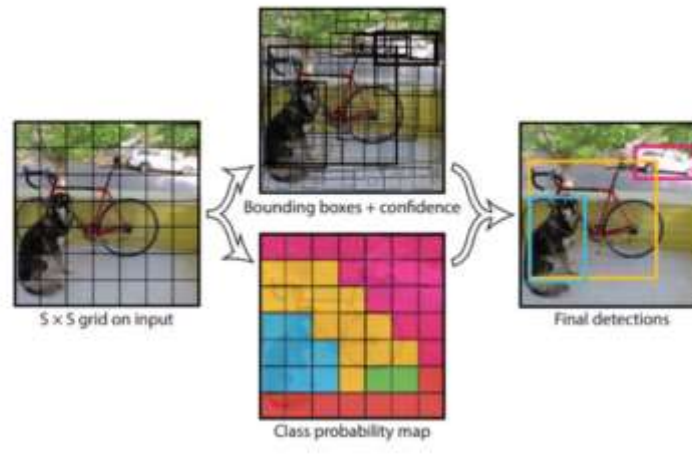
Mask	CNN Architecture	Metrics (%)			
		P	R	IoU	F
Schroeder et al. (non-CNN)	-	88.1	70.2	64.1	78.1
	U-Net (10c)	86.9	88.4	78.0	87.7
	U-Net (3c)	89.5	88.5	73.6	84.5
	U-Net-Light (3c)	89.0	78.8	71.5	83.6
Murphy et al. (non-CNN)	-	76.6	96.1	74.3	85.2
	U-Net (10c)	74.8	96.9	73.0	84.4
	U-Net (3c)	72.7	97.2	71.2	83.2
	U-Net-Light (3c)	75.5	96.9	73.7	84.9
Kumar-Roy (non-CNN)	-	82.3	68.4	59.7	74.7
	U-Net (10c)	82.2	94.2	78.3	87.8
	U-Net (3c)	84.5	93.4	79.8	88.8
	U-Net-Light (3c)	78.8	96.9	76.9	86.9
Intersection (non-CNN)	-	92.6	57.6	55.1	71.0
Intersection {	Schroeder et al.	91.8	75.4	70.6	82.8
	Murphy et al.	90.8	73.5	68.4	81.2
	Kumar - Roy	90.8	73.8	67.8	80.0
Voting (non-CNN)	-	87.7	79.1	71.2	83.2
Voting {	Schroeder et al.	83.6	94.0	79.3	88.5
	Murphy et al.	86.4	93.0	81.1	89.6
	Kumar - Roy	87.2	92.4	81.4	89.7

Object detection

Faster R-CNN



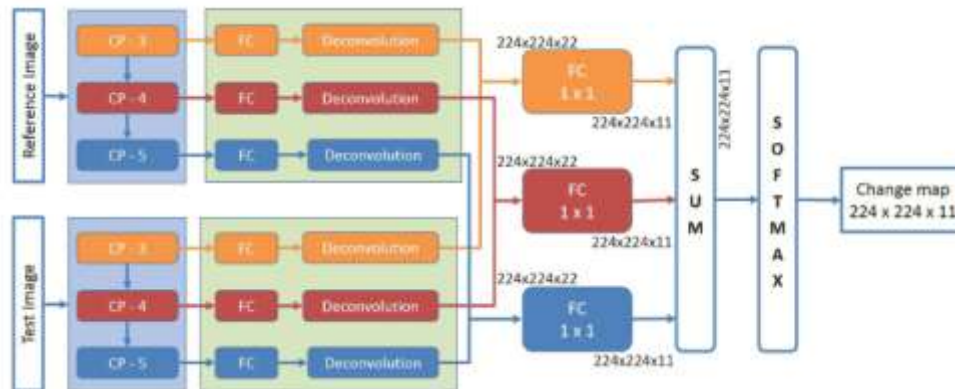
YOLO — You Only Look Once



Change detection

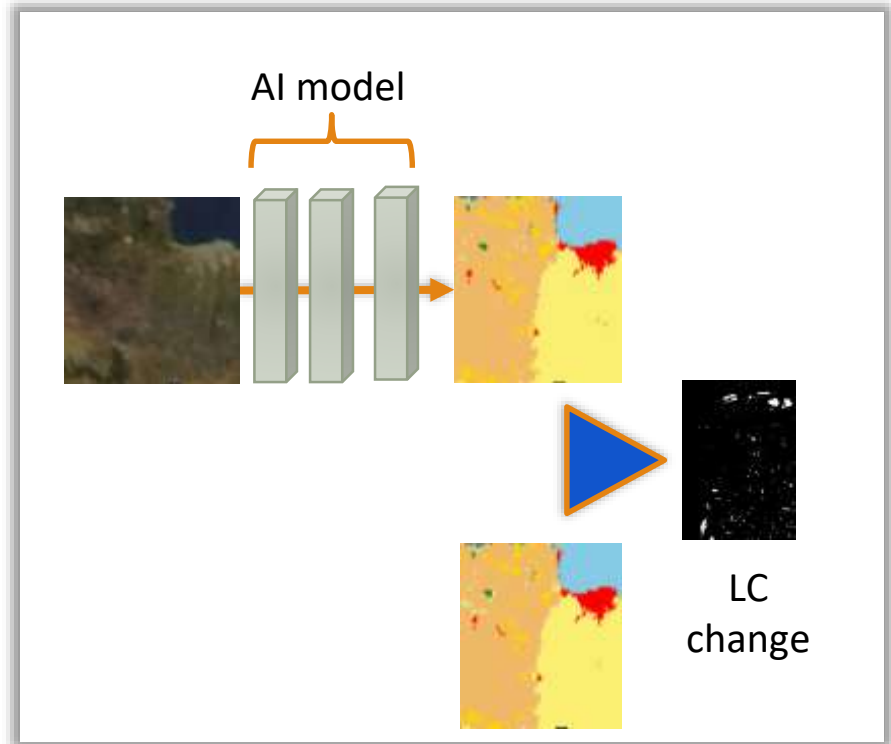
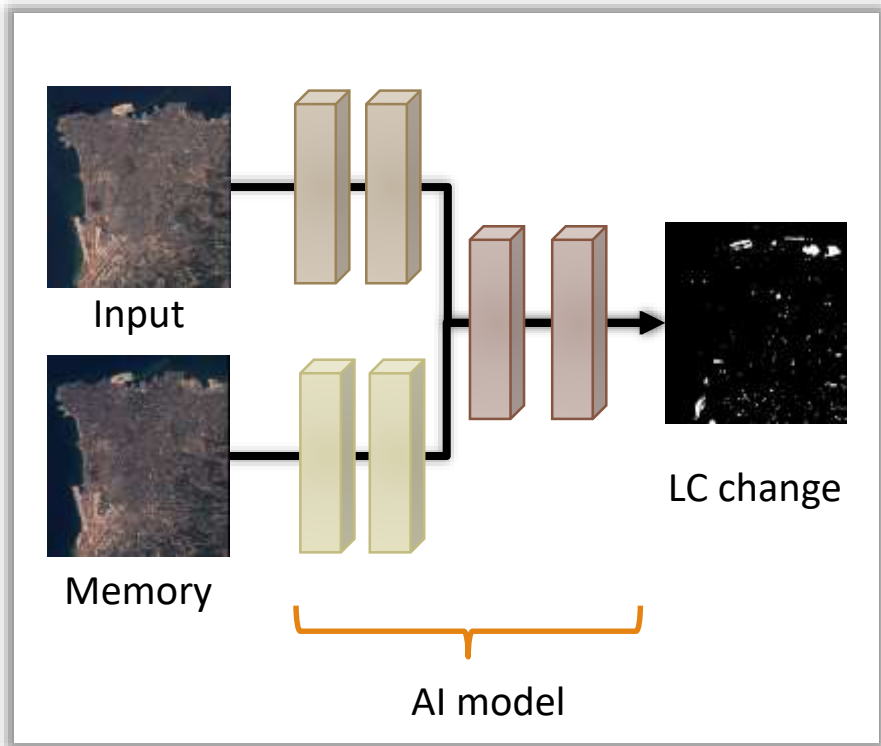
Challenges:

- Changes in lighting, atmospheric conditions, and seasonality
- Lack of High-Quality, Annotated Data
- Time-series dynamics
- Heterogeneity of Data



Varghese, Ashley, et al. "ChangeNet: A deep learning architecture for visual change detection." *Proceedings of the European conference on computer vision (ECCV) workshops*. 2018.

Land cover change detection



Performance metrics: memory, complexity, speed, accuracy

Performance

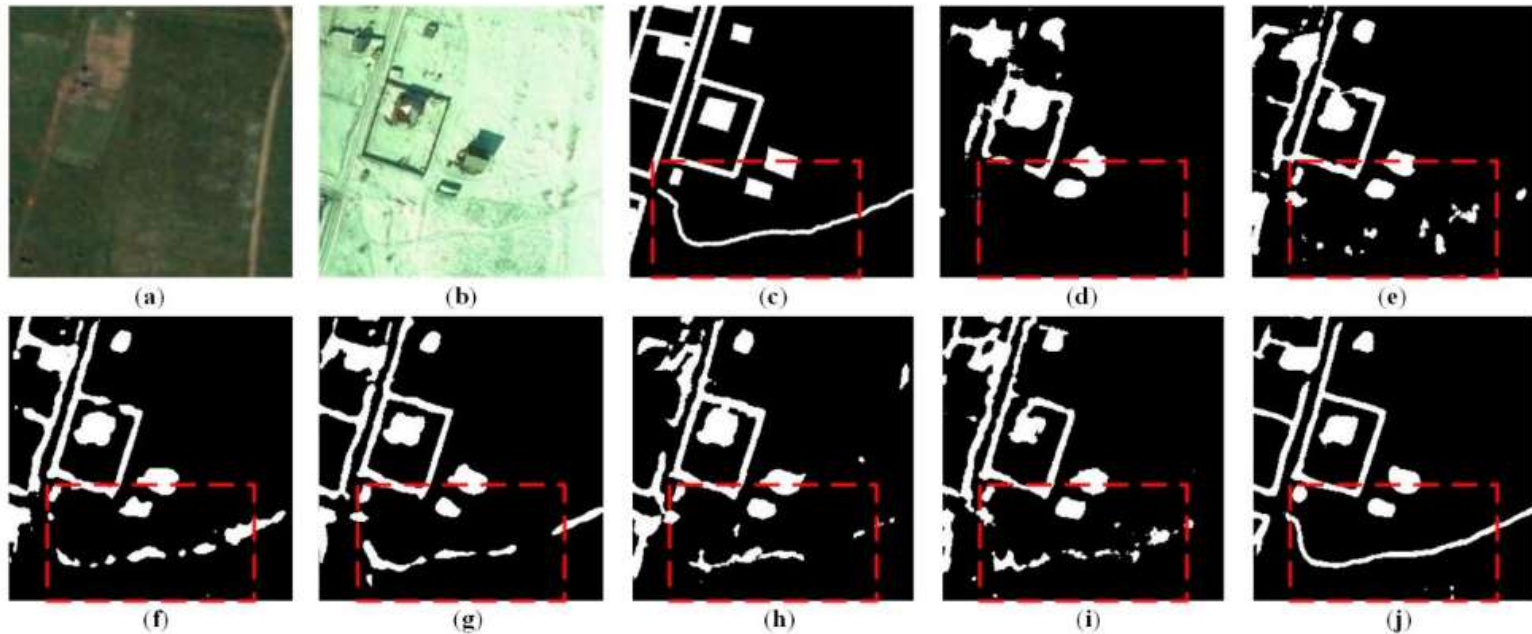


Figure 5. Visual comparison of CD results using various DL methods for area 6: (a) image T1, (b) image T2, (c) reference change map, (d) CDNet, (e) FC-EF, (f) FC-Siam-conc, (g) FC-Siam-diff, (h) FC-EF-Res, (i) FCN-PP, and (j) U-Net++. The changed parts are marked in white while the unchanged are in black.

Shafique, Ayesha, et al. "Deep learning-based change detection in remote sensing images: A review." *Remote Sensing* 14.4 (2022): 871.

TensorFlow

Deep learning library, open-sourced by Google
(11/2015)

TensorFlow provides primitives for

- defining functions on tensors
- automatically computing their derivatives



What is a tensor

What is a computational graph

Material from lecture by Bharath Ramsundar, March 2018, Stanford

Introduction to Keras

Official high-level API of TensorFlow

- Python
- 250K developers

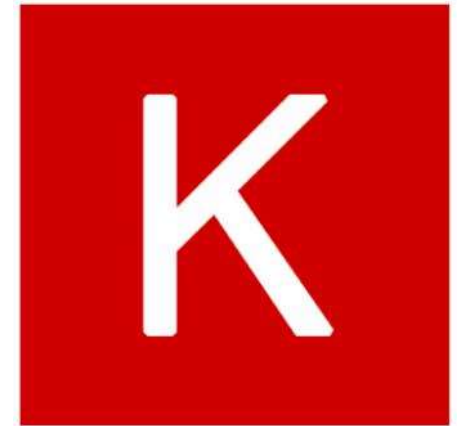
Same front-end <-> Different back-ends

- TensorFlow (Google)
- CNTK (Microsoft)
- MXNet (Apache)
- Theano (RIP)

Hardware

- GPU (Nvidia)
- CPU (Intel/AMD)
- TPU (Google)

Companies: Netflix, Uber, Google, Nvidia...



Material from lecture by Francois Chollet, 2018, Stanford

Keras models

Installation

- Anaconda -> Tensorflow -> Keras

Build-in

- Conv1D, Conv2D, Conv3D...
- MaxPooling1D, MaxPooling2D, MaxPooling3D...
- Dense, Activation, RNN...

The Sequential Model

- Very simple
- Single-input, Single-output, sequential layer stacks

The functional API

- Mix & Match
- Multi-input, multi-output, arbitrary static graph topologies

Sequential

```
>> from keras.models import Sequential  
>> model = Sequential()  
>> from keras.layers import Dense  
>> model.add(Dense(units=64, activation='relu', input_dim=100))  
>> model.add(Dense(units=10, activation='softmax'))  
>> model.compile(loss='categorical_crossentropy',  
optimizer='sgd', metrics=['accuracy'])  
>> model.fit(x_train, y_train, epochs=5, batch_size=32)  
>> loss_and_metrics = model.evaluate(x_test, y_test,  
batch_size=128)  
>> classes = model.predict(x_test)
```


Functional

```
>> from keras.layers import Input, Dense
>> from keras.models import Model
>> inputs = Input(shape=(784,))
>> x = Dense(64, activation='relu')(inputs)
>> x = Dense(64, activation='relu')(x)
>> predictions = Dense(10, activation='softmax')(x)
>> model = Model(inputs=inputs, outputs=predictions)
>> model.compile(optimizer='rmsprop',
loss='categorical_crossentropy', metrics=['accuracy'])
>> model.fit(data, labels)
```

TensorFlow

Written in C++ and is, as a result, very fast and efficient.

Feature rich; TensorFlow can be used for training data as well as for inference.

Very good documentation; TensorFlow has many users and a big community which has led to strong documentation.

High popularity; TensorFlow has established itself as the most used ML library over a number of years now.

Many APIs available; TensorFlow is a library with a rich choice of easy to use APIs.

Supports JavaScript: TensorFlow supports JavaScript, C++ and Java in addition to Python.

For Mobile & IoT, inferences can be performed with TensorFlow Lite on mobile devices such as Android or iOS, as well as on Edge TPU or Raspberry Pi.

PyTorch

Written in Python making it more accessible and flattening the learning curve. However, the C++ core means PyTorch is still quite fast.

Very flexible; as data size can also be changed during data training.

Popular at research level; PyTorch was by far the most talked about ML library at CVPR, one of the most important computer vision conferences.

Rapid growth in popularity in both business and research use cases.

Many libraries available; PyTorch is composed of multiple libraries and platforms.

Python-based; PyTorch allows developers to write code in Python

PyTorch API; the PyTorch API is often preferred as it is better designed - plus TensorFlow has historically changed their API frequently.

Hands-on

